



GAN: Projects

Contents:

I. Project I: Colorization

- A. Basic Structure
- B. Updated Version
- C. GAN Based Method
- D. Tricky Losses: Perceptual Loss + TV Loss

II. Project II: Face Frontalization

- E. Modified Backbone
- F. Tricky Losses: Symmetrical Loss + Extra Loss

Contents:

III. Project I: Revisit

G. Temporal Consistent

IV. Project II: Revisit

H. Pair with Unpair

I. Project I: Colorization



I. Project I: Colorization

A. Basic Structure

- Straight-line Structure

```
class Net(nn.Module):  
    def __init__(self, input_size=256):  
        super(Net, self).__init__()  
        # ResNet - First layer accepts grayscale images,  
        # and we take only the first few layers of ResNet for this task  
        resnet = models.resnet18(num_classes=100)  
        resnet.conv1.weight = nn.Parameter(resnet.conv1.weight.sum(dim=1).unsqueeze(1))  
        self.midlevel_resnet = nn.Sequential(*list(resnet.children())[0:6])  
        RESNET_FEATURE_SIZE = 128
```

```
class Net(nn.Module):  
    def __init__(self, input_size=256):  
        super(Net, self).__init__()  
        # ResNet - First layer accepts grayscale images,  
        # and we take only the first few layers of ResNet for this task  
        resnet = models.resnet18(num_classes=100)  
        resnet.conv1.weight = nn.Parameter(resnet.conv1.weight.sum(dim=1).unsqueeze(1))  
        self.midlevel_resnet = nn.Sequential(*list(resnet.children())[0:6])  
        RESNET_FEATURE_SIZE = 128  
        # Upsampling Network  
        self.upsample = nn.Sequential(  
            nn.Conv2d(RESNET_FEATURE_SIZE, 128, kernel_size=3, stride=1, padding=1),  
            .....,  
            nn.Conv2d(32, 3, kernel_size=3, stride=1, padding=1),  
            nn.Upsample(scale_factor=2)  
        )  
  
    def forward(self, input):  
        midlevel_features = self.midlevel_resnet(input)  
        output = self.upsample(midlevel_features)  
        return output
```

I. Project I: Colorization

A. Basic Structure

- **Straight-line Structure**

```
class Net(nn.Module):
    def __init__(self, input_size=256):
        super(Net, self).__init__()
        # ResNet - First layer accepts grayscale images,
        # and we take only the first few layers of ResNet for this task
        resnet = models.resnet18(num_classes=100)
        resnet.conv1.weight = nn.Parameter(resnet.conv1.weight.sum(dim=1).unsqueeze(1))
        self.midlevel_resnet = nn.Sequential(*list(resnet.children())[0:6])
        RESNET_FEATURE_SIZE = 128
```

nn.Sequential():

It contains sequential/ordered executions

sum all parameters along 1st dim
[64, 3, 7, 7] → [64, 7, 7]

```
>>> print(hei[0])
Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
>>> print(hei[1])
BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
>>> print(hei[2])
ReLU(inplace=True)
>>> print(hei[3])
MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
```

expand 1 dim at 1st spot
[64, 7, 7] → [64, 1, 7, 7]

Set conv1.weight from 3-dim
to 1-dim [gray image]

```
>>> print(hei[4])
Sequential(
    (0): BasicBlock(
        (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
        (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
        (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
    (1): BasicBlock(
        (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
        (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
        (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
)
```

unpack a list

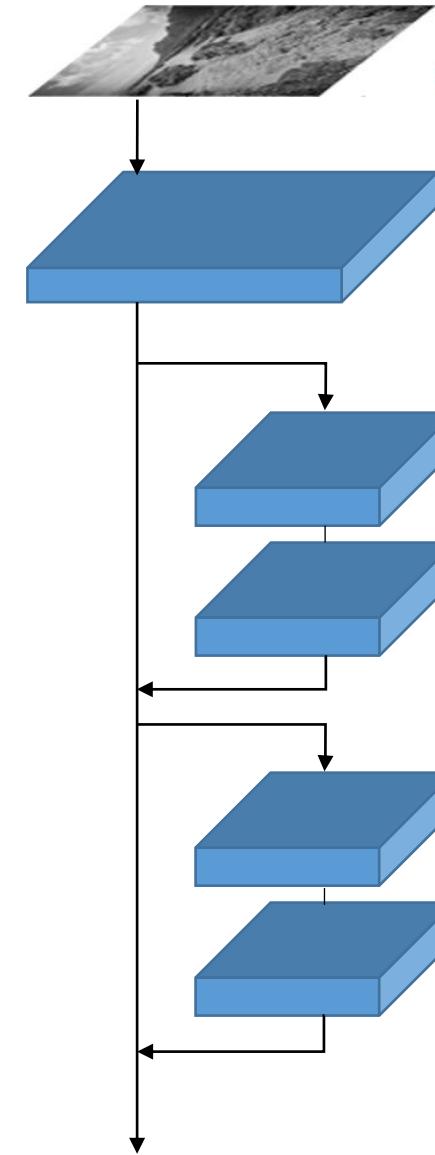
resnet layers/operations

I. Project I: Colorization

A. Basic Structure

- Straight-line Structure

```
class Net(nn.Module):  
    def __init__(self, input_size=256):  
        super(Net, self).__init__()  
        # ResNet - First layer accepts grayscale images,  
        # and we take only the first few layers of ResNet for this task  
        resnet = models.resnet18(num_classes=100)  
        resnet.conv1.weight = nn.Parameter(resnet.conv1.weight.sum(dim=1).unsqueeze(1))  
        self.midlevel_resnet = nn.Sequential(*list(resnet.children())[0:6])  
        RESNET_FEATURE_SIZE = 100  
        >>> print(hei[0])  
        Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)  
        >>> print(hei[1])  
        BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
        >>> print(hei[2])  
        ReLU(inplace=True)  
        >>> print(hei[3])  
        MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)  
        >>> print(hei[4])  
        Sequential(  
            (0): BasicBlock(  
                (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)  
                (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
                (relu): ReLU(inplace=True)  
                (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)  
                (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
            )  
            (1): BasicBlock(  
                (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)  
                (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
                (relu): ReLU(inplace=True)  
                (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)  
                (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
            )  
        )
```



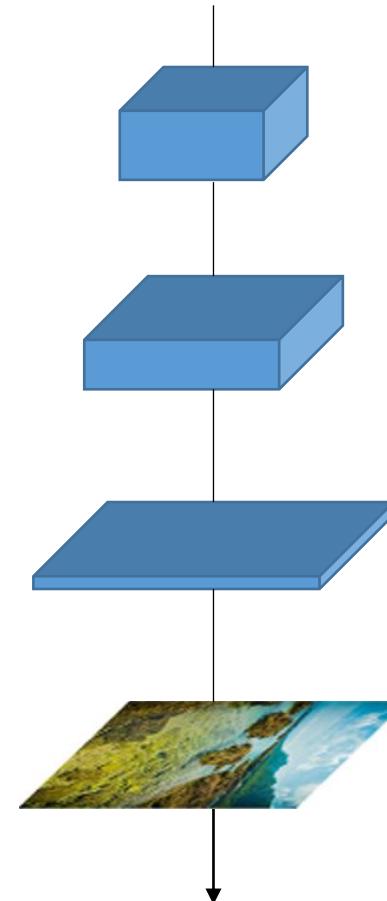
I. Project I: Colorization

A. Basic Structure

- **Straight-line Structure**

```
class Net(nn.Module):
    def __init__(self, input_size=256):
        .....
        RESNET_FEATURE_SIZE = 128
        # Upsampling Network
        self.upsample = nn.Sequential(
            nn.Conv2d(RESNET_FEATURE_SIZE, 128, kernel_size=3, stride=1, padding=1),
            nn.BatchNorm2d(128),
            nn.ReLU(),
            nn.Upsample(scale_factor=2),
            nn.Conv2d(128, 64, kernel_size=3, stride=1, padding=1),
            nn.BatchNorm2d(64),
            nn.ReLU(),
            nn.Conv2d(64, 64, kernel_size=3, stride=1, padding=1),
            nn.BatchNorm2d(64),
            nn.ReLU(),
            nn.Upsample(scale_factor=2),
            nn.Conv2d(64, 32, kernel_size=3, stride=1, padding=1),
            nn.BatchNorm2d(32),
            nn.ReLU(),
            nn.Conv2d(32, 3, kernel_size=3, stride=1, padding=1),
            nn.Upsample(scale_factor=2)
        )

    def forward(self, input):
        midlevel_features = self.midlevel_resnet(input)
        output = self.upsample(midlevel_features)
        return output
```



I. Project I: Colorization

B. Updated Version

- **UNet Structure**

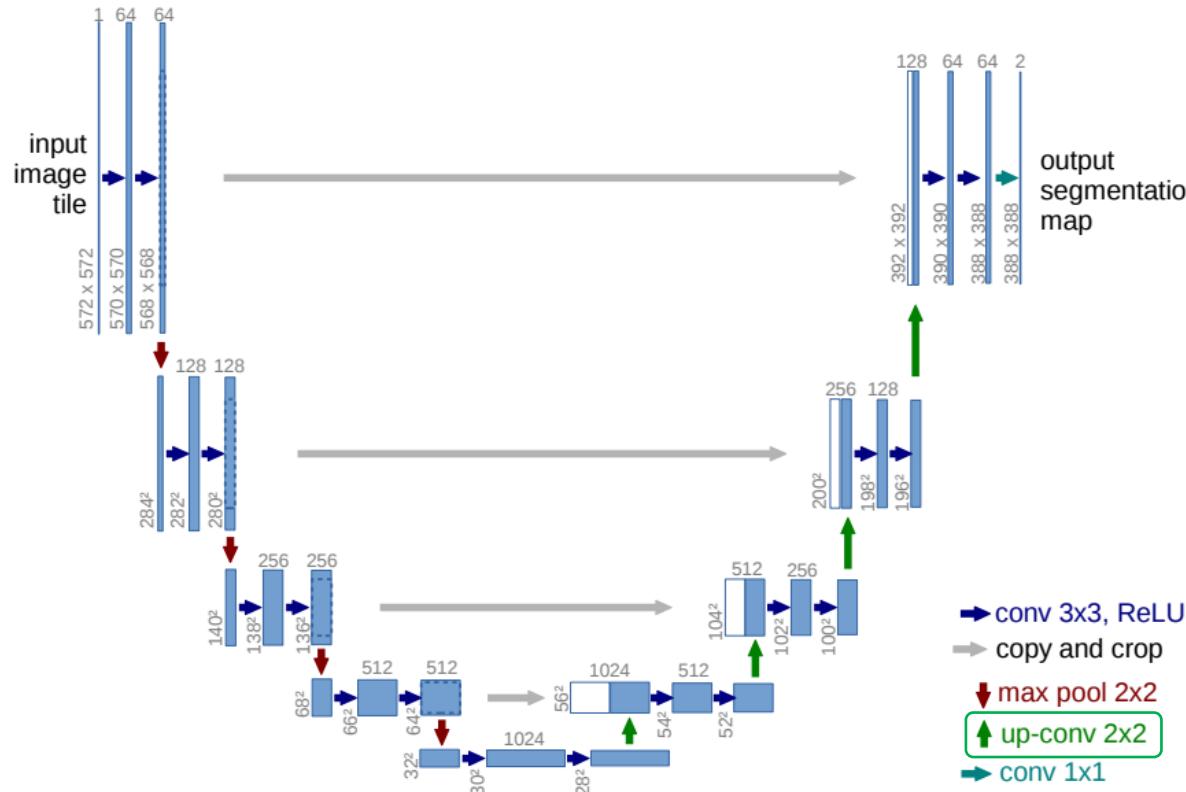
```
class Net(nn.Module):
    def __init__(self, input_size=256):
        .....
        RESNET_FEATURE_SIZE = 128
        # Upsampling Network
        self.upsample = nn.Sequential(
            nn.Conv2d(RESNET_FEATURE_SIZE, 128, kernel_size=3, stride=1, padding=1),
            nn.BatchNorm2d(128),
            nn.ReLU(),
            nn.Upsample(scale_factor=2),
            nn.Conv2d(128, 64, kernel_size=3, stride=1, padding=1),
            nn.BatchNorm2d(64),
            nn.ReLU(),
            nn.Conv2d(64, 64, kernel_size=3, stride=1, padding=1),
            nn.BatchNorm2d(64),
            nn.ReLU(),
            nn.Upsample(scale_factor=2),
            nn.Conv2d(64, 32, kernel_size=3, stride=1, padding=1),
            nn.BatchNorm2d(32),
            nn.ReLU(),
            nn.Conv2d(32, 3, kernel_size=3, stride=1, padding=1),
            nn.Upsample(scale_factor=2)
        )

        def forward(self, input):
            midlevel_features = self.midlevel_resnet(input) Encoder
            output = self.upsample(midlevel_features) Decoder
            return output
```

I. Project I: Colorization

B. Updated Version

- **UNet Structure [2015, Olaf Ronneberger]**



I. Project I: Colorization

B. Updated Version

- **Deconv / Transpose Conv / Strided Conv, Upsample**

Normal convolution

$$\begin{pmatrix} a & b & c \\ d & e & f \\ g & h & i \end{pmatrix} \otimes \begin{pmatrix} 1 & 2 \\ 3 & 4 \end{pmatrix} = \begin{pmatrix} 1a + 2b + 3d + 4e & 1b + 2c + 3e + 4f \\ 1d + 2e + 3g + 4h & 1e + 2f + 3h + 4i \end{pmatrix} \equiv \begin{pmatrix} A & B \\ C & D \end{pmatrix}$$

I. Project I: Colorization

B. Updated Version

- **Deconv / Transpose Conv / Strided Conv, Upsample**

Normal convolution

$$\begin{pmatrix} a & b & c \\ d & e & f \\ g & h & i \end{pmatrix} \otimes \begin{pmatrix} 1 & 2 \\ 3 & 4 \end{pmatrix} = \begin{pmatrix} 1a + 2b + 3d + 4e & 1b + 2c + 3e + 4f \\ 1d + 2e + 3g + 4h & 1e + 2f + 3h + 4i \end{pmatrix} \equiv \begin{pmatrix} A & B \\ C & D \end{pmatrix} \quad 3 \times 3 \rightarrow 2 \times 2$$

Normal convolution in matrix multiplication

$$\begin{pmatrix} A \\ B \\ C \\ D \end{pmatrix} = \begin{pmatrix} 1 & 2 & 0 & 3 & 4 & 0 & 0 & 0 & 0 \\ 0 & 1 & 2 & 0 & 3 & 4 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 2 & 0 & 3 & 4 & 0 \\ 0 & 0 & 0 & 0 & 1 & 2 & 0 & 3 & 4 \end{pmatrix} * \begin{pmatrix} a \\ b \\ c \\ d \\ e \\ f \\ g \\ h \\ i \end{pmatrix}$$

变小了

$9 \rightarrow 4$

I. Project I: Colorization

B. Updated Version

- Deconv / Transpose Conv / Strided Conv, Upsample

Normal convolution in matrix multiplication

$$\begin{pmatrix} A \\ B \\ C \\ D \end{pmatrix} = \begin{pmatrix} 1 & 2 & 0 & 3 & 4 & 0 & 0 & 0 & 0 \\ 0 & 1 & 2 & 0 & 3 & 4 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 2 & 0 & 3 & 4 & 0 \\ 0 & 0 & 0 & 0 & 1 & 2 & 0 & 3 & 4 \end{pmatrix} * \begin{pmatrix} a \\ b \\ c \\ d \\ e \\ f \\ g \\ h \end{pmatrix}$$

变小了

$9 \rightarrow 4$

$\begin{pmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \\ x_5 \\ x_6 \\ x_7 \\ x_8 \\ x_9 \end{pmatrix}$ Normal transpose convolution in matrix multiplication

$$\begin{pmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \\ x_5 \\ x_6 \\ x_7 \\ x_8 \\ x_9 \end{pmatrix} = \begin{pmatrix} 1 & 2 & 0 & 3 & 4 & 0 & 0 & 0 & 0 \\ 0 & 1 & 2 & 0 & 3 & 4 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 2 & 0 & 3 & 4 & 0 \\ 0 & 0 & 0 & 0 & 1 & 2 & 0 & 3 & 4 \end{pmatrix}^T * \begin{pmatrix} A' \\ B' \\ C' \\ D' \end{pmatrix} = \begin{pmatrix} 1 & 0 & 0 & 0 \\ 2 & 1 & 0 & 0 \\ 0 & 2 & 0 & 0 \\ 3 & 0 & 1 & 0 \\ 4 & 3 & 2 & 1 \\ 0 & 4 & 0 & 2 \\ 0 & 0 & 3 & 0 \\ 0 & 0 & 4 & 3 \\ 0 & 0 & 0 & 4 \end{pmatrix} * \begin{pmatrix} A' \\ B' \\ C' \\ D' \end{pmatrix}$$

变大了

$4 \rightarrow 9$

I. Project I: Colorization

B. Updated Version

- Deconv / Transpose Conv / Strided Conv, Upsample

Normal transpose convolution in matrix multiplication

$$\begin{pmatrix} x_1 & x_2 & x_3 \\ x_4 & x_5 & x_6 \\ x_7 & x_8 & x_9 \end{pmatrix} \equiv \begin{pmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \\ x_5 \\ x_6 \\ x_7 \\ x_8 \\ x_9 \end{pmatrix} = \begin{pmatrix} 1 & 2 & 0 & 3 & 4 & 0 & 0 & 0 & 0 \\ 0 & 1 & 2 & 0 & 3 & 4 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 2 & 0 & 3 & 4 & 0 \\ 0 & 0 & 0 & 0 & 1 & 2 & 0 & 3 & 4 \end{pmatrix}^T * \begin{pmatrix} A' \\ B' \\ C' \\ D' \end{pmatrix} = \begin{pmatrix} 1 & 0 & 0 & 0 \\ 2 & 1 & 0 & 0 \\ 0 & 2 & 0 & 0 \\ 3 & 0 & 1 & 0 \\ 4 & 3 & 2 & 1 \\ 0 & 4 & 0 & 2 \\ 0 & 0 & 3 & 0 \\ 0 & 0 & 4 & 3 \\ 0 & 0 & 0 & 4 \end{pmatrix} * \begin{pmatrix} A' \\ B' \\ C' \\ D' \end{pmatrix} \equiv \begin{pmatrix} A' & B' \\ C' & D' \end{pmatrix}$$

变大了
 $4 \rightarrow 9$

By transposing the kernel matrix,
we can upscale the output vector comparing to the size of the input vector

I. Project I: Colorization

B. Updated Version

- **Deconv / Transpose Conv / Strided Conv, Upsample**

Normal **transpose** convolution in matrix multiplication

$$\begin{pmatrix} x_1 & x_2 & x_3 \\ x_4 & x_5 & x_6 \\ x_7 & x_8 & x_9 \end{pmatrix} \equiv \begin{pmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \\ x_5 \\ x_6 \\ x_7 \\ x_8 \\ x_9 \end{pmatrix} = \begin{pmatrix} 1 & 2 & 0 & 3 & 4 & 0 & 0 & 0 & 0 \\ 0 & 1 & 2 & 0 & 3 & 4 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 2 & 0 & 3 & 4 & 0 \\ 0 & 0 & 0 & 0 & 1 & 2 & 0 & 3 & 4 \end{pmatrix}^T * \begin{pmatrix} A' \\ B' \\ C' \\ D' \end{pmatrix} = \begin{pmatrix} 1 & 0 & 0 & 0 \\ 2 & 1 & 0 & 0 \\ 0 & 2 & 0 & 0 \\ 3 & 0 & 1 & 0 \\ 4 & 3 & 2 & 1 \\ 0 & 4 & 0 & 2 \\ 0 & 0 & 3 & 0 \\ 0 & 0 & 4 & 3 \\ 0 & 0 & 0 & 4 \end{pmatrix} * \begin{pmatrix} A' \\ B' \\ C' \\ D' \end{pmatrix} = \begin{pmatrix} A' \\ 2A' + B' \\ 2B' \\ 3A' + C' \\ 4A' + 3B' + 2B' + C' \\ 4B' + 2D' \\ 3C' \\ 4C' + 3D' \\ 4D' \end{pmatrix}$$

I. Project I: Colorization

B. Updated Version

- Deconv / Transpose Conv / Strided Conv, Upsample

$$\begin{pmatrix} A' & B' \\ C' & D' \end{pmatrix} \otimes \begin{pmatrix} 1 & 2 \\ 3 & 4 \end{pmatrix} = \begin{pmatrix} A' & B' & 0 \\ C' & D' & 0 \\ 0 & 0 & 0 \end{pmatrix} + \begin{pmatrix} 0 & 2A' & 2B' \\ 0 & 2C' & 2D' \\ 0 & 0 & 0 \end{pmatrix} + \begin{pmatrix} 0 & 0 & 0 \\ 3A' & 3B' & 0 \\ 3C' & 3D' & 0 \end{pmatrix} + \begin{pmatrix} 0 & 0 & 0 \\ 0 & 4A' & 4B' \\ 0 & 4C' & 4D' \end{pmatrix}$$

$$= \begin{pmatrix} A' & 2A' + B' & 2B' \\ 3A' + C' & 4A' + 3B' + 2B' + C' & 4B' + 2D' \\ 3C' & 4C' + 3D' & 4D' \end{pmatrix} \equiv \begin{pmatrix} A' \\ 2A' + B' \\ 2B' \\ 3A' + C' \\ 4A' + 3B' + 2B' + C' \\ 4B' + 2D' \\ 3C' \\ 4C' + 3D' \\ 4D' \end{pmatrix}$$

变大了

$2 \times 2 \rightarrow 3 \times 3$

I. Project I: Colorization

B. Updated Version

- Deconv / Transpose Conv / Strided Conv, Upsample

$$\begin{pmatrix} 1 & 2 \\ 3 & 4 \end{pmatrix} \otimes \begin{pmatrix} A' & B' \\ C' & D' \end{pmatrix} = \begin{pmatrix} A' & 2A' & 0 \\ 3A' & 4A' & 0 \\ 0 & 0 & 0 \end{pmatrix} + \begin{pmatrix} 0 & B' & 2B' \\ 0 & 3B' & 4B' \\ 0 & 0 & 0 \end{pmatrix} + \begin{pmatrix} 0 & 0 & 0 \\ C' & 2C' & 0 \\ 3C' & 4C' & 0 \end{pmatrix} + \begin{pmatrix} 0 & 0 & 0 \\ 0 & D' & 2D' \\ 0 & 3D' & 4D' \end{pmatrix}$$

$$= \begin{pmatrix} A' & 2A' + B' & 2B' \\ 3A' + C' & 4A' + 3B' + 2B' + C' & 4B' + 2D' \\ 3C' & 4C' + 3D' & 4D' \end{pmatrix} \equiv \begin{pmatrix} A' \\ 2A' + B' \\ 2B' \\ 3A' + C' \\ 4A' + 3B' + 2B' + C' \\ 4B' + 2D' \\ 3C' \\ 4C' + 3D' \\ 4D' \end{pmatrix}$$

变大了
 $2 \times 2 \rightarrow 3 \times 3$

I. Project I: Colorization

B. Updated Version

- **Deconv / Transpose Conv / Strided Conv, Upsample**

```
>>> input = torch.arange(1, 5, dtype=torch.float32).view(1, 1, 2, 2)
>>> input
tensor([[[[ 1.,  2.],
          [ 3.,  4.]]]])
```

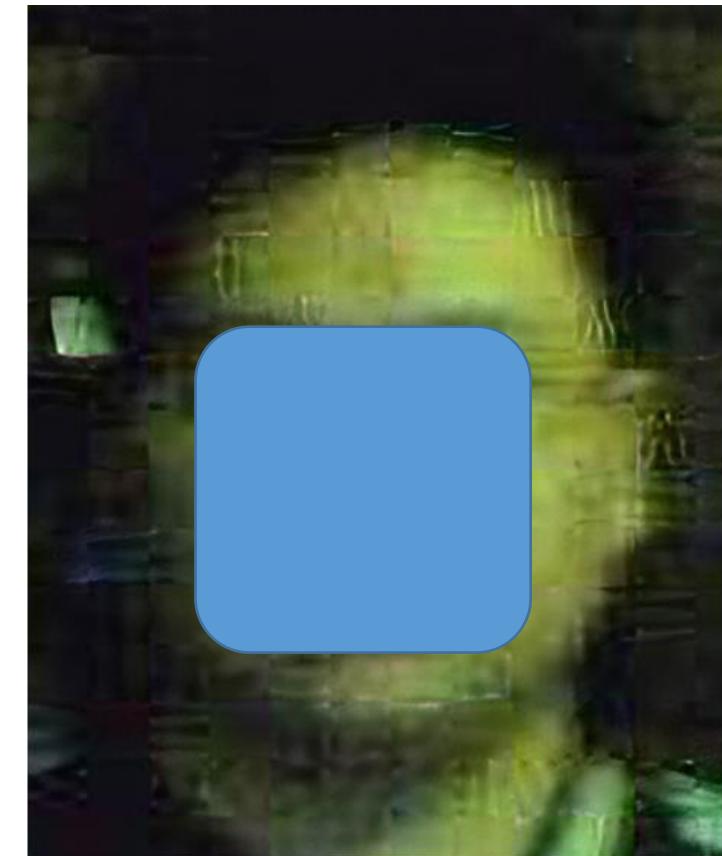
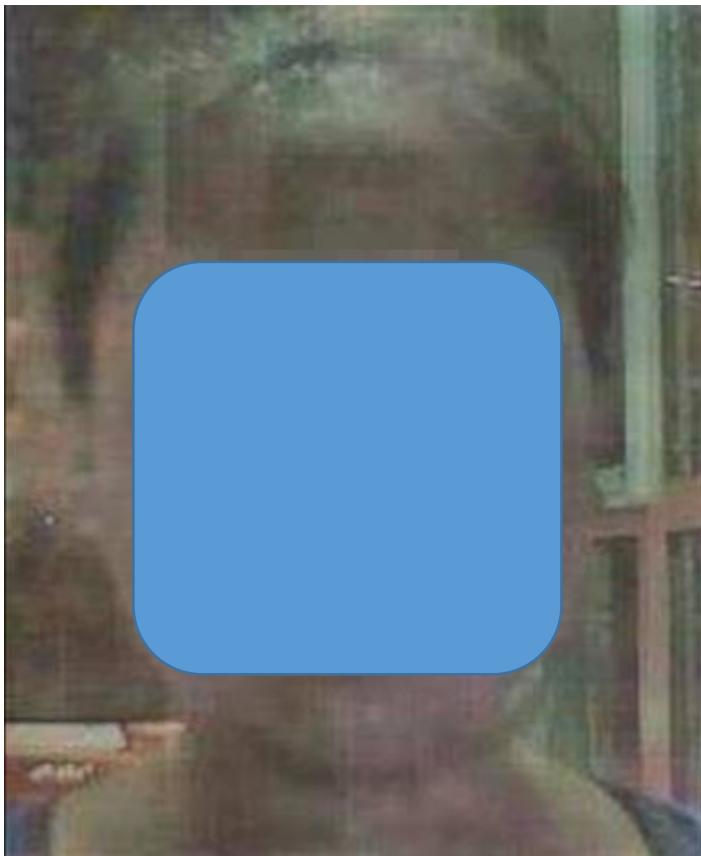


```
>>> m = nn.Upsample(scale_factor=2, mode='nearest')      >>> m = nn.Upsample(scale_factor=2, mode='bilinear')
>>> m(input)                                         >>> m(input)
tensor([[[[ 1.,  1.,  2.,  2.],
          [ 1.,  1.,  2.,  2.],
          [ 3.,  3.,  4.,  4.],
          [ 3.,  3.,  4.,  4.]]]])                         tensor([[[[ 1.0000,  1.2500,  1.7500,  2.0000],
          [ 1.5000,  1.7500,  2.2500,  2.5000],
          [ 2.5000,  2.7500,  3.2500,  3.5000],
          [ 3.0000,  3.2500,  3.7500,  4.0000]]]])
```

I. Project I: Colorization

B. Updated Version

- Transpose Conv / Strided Conv bad effect – **chessboard pattern**



I. Project I: Colorization

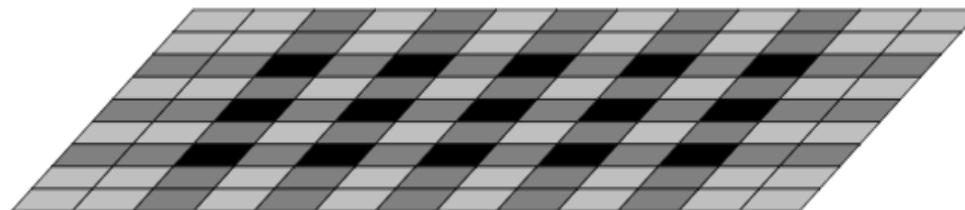
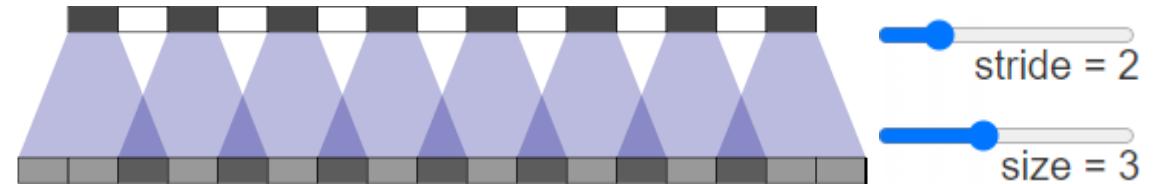
B. Updated Version

- Transpose Conv / Strided Conv bad effect – **chessboard pattern**

Reason: [great illustration](#)

How to solve:

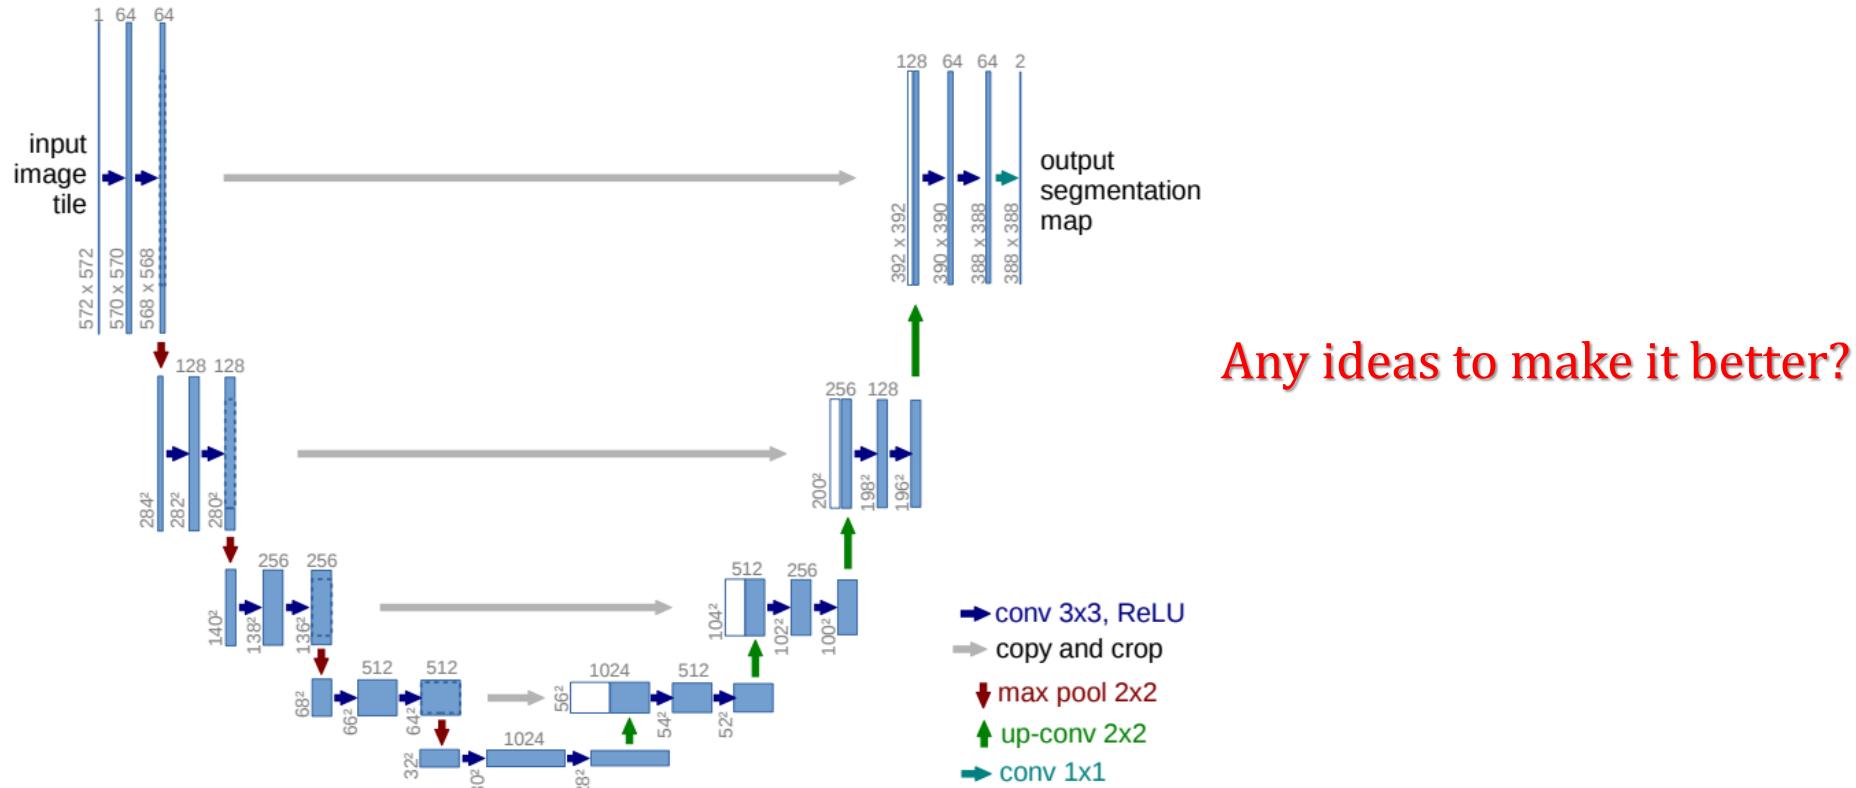
1. Ticks provided in the link
[sophisticated design] → deprecated
2. Several training strategies,
[we'll talk later](#)

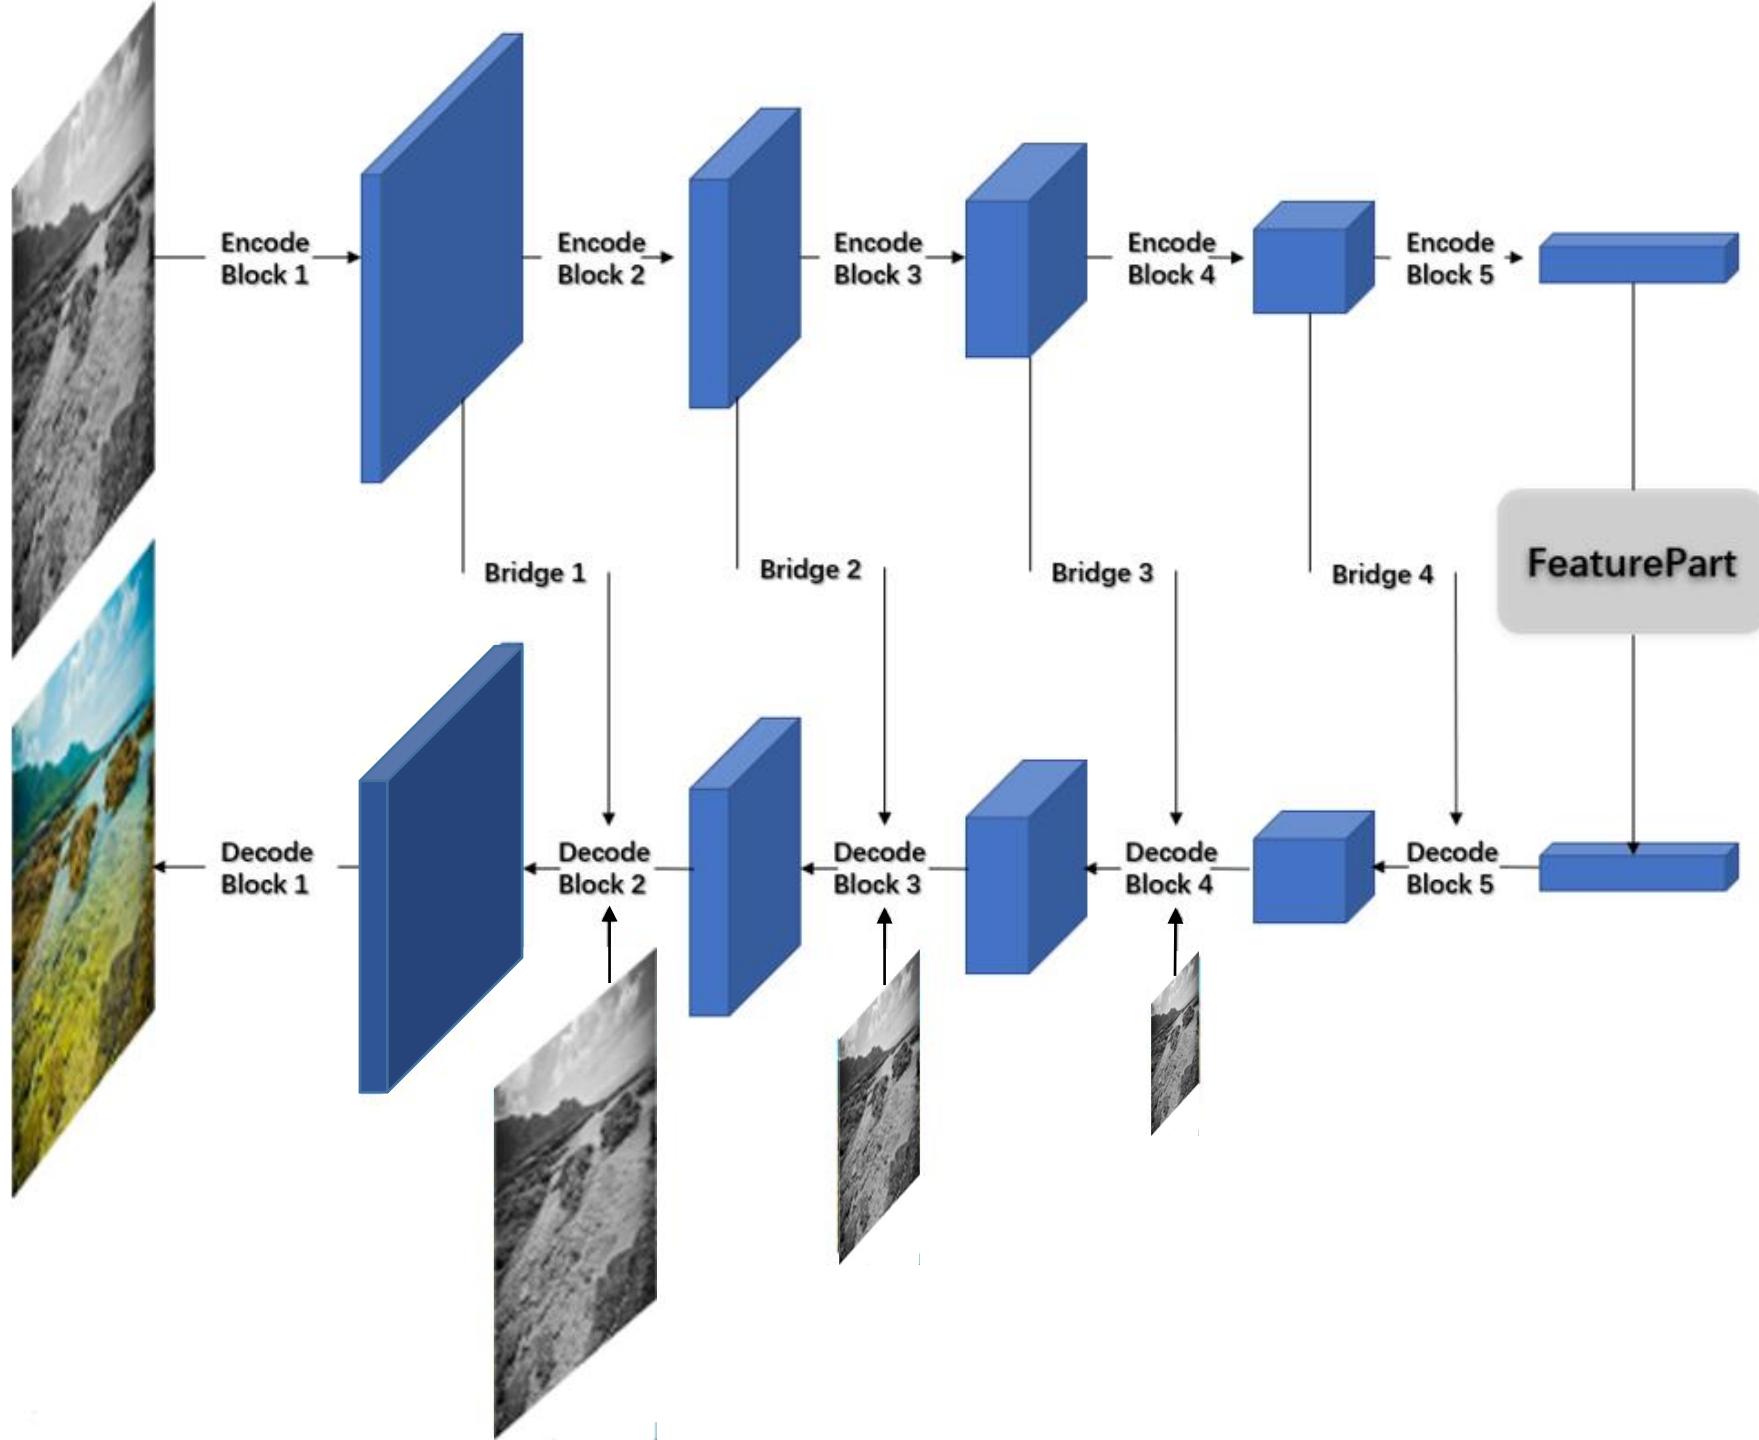


I. Project I: Colorization

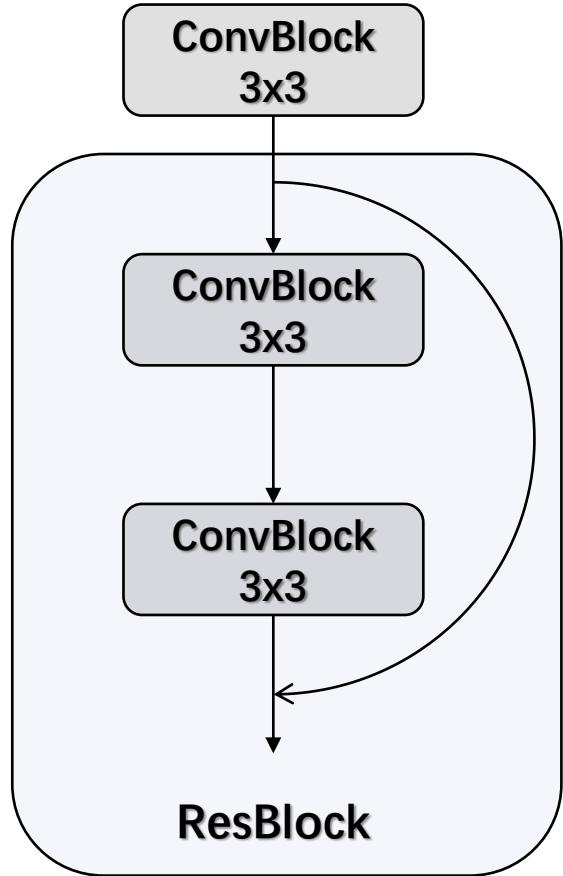
B. Updated Version

- **UNet Structure [2015, Olaf Ronneberger]**

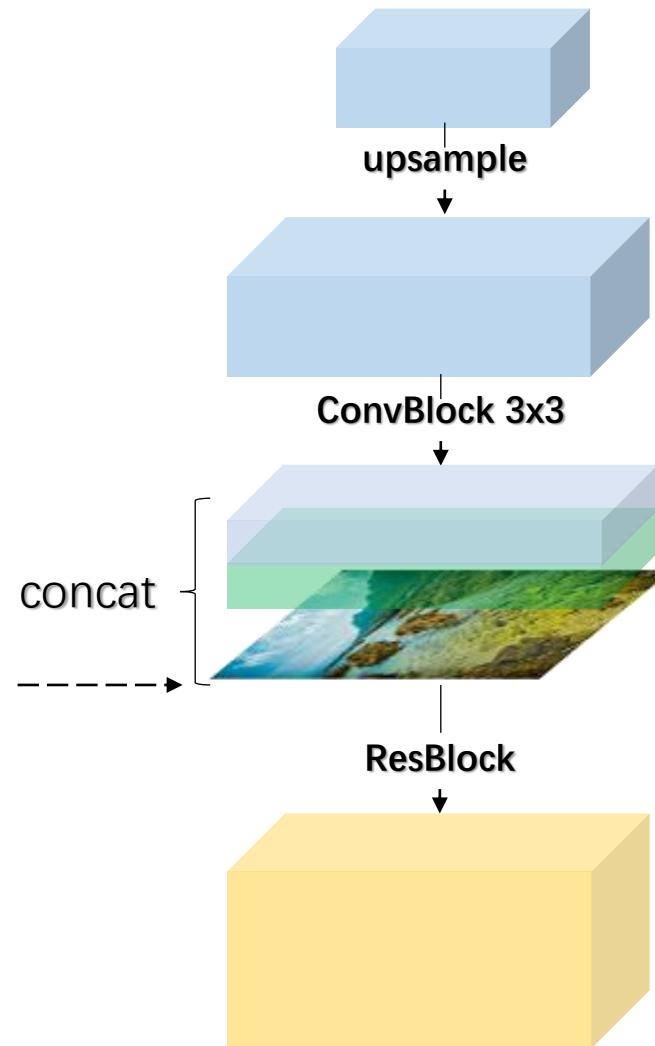




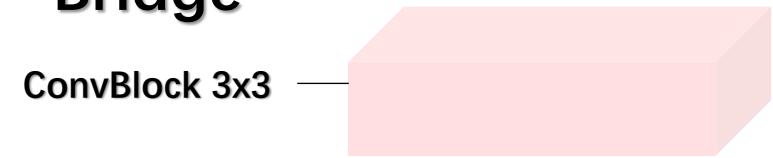
Encode Block



Decode Block



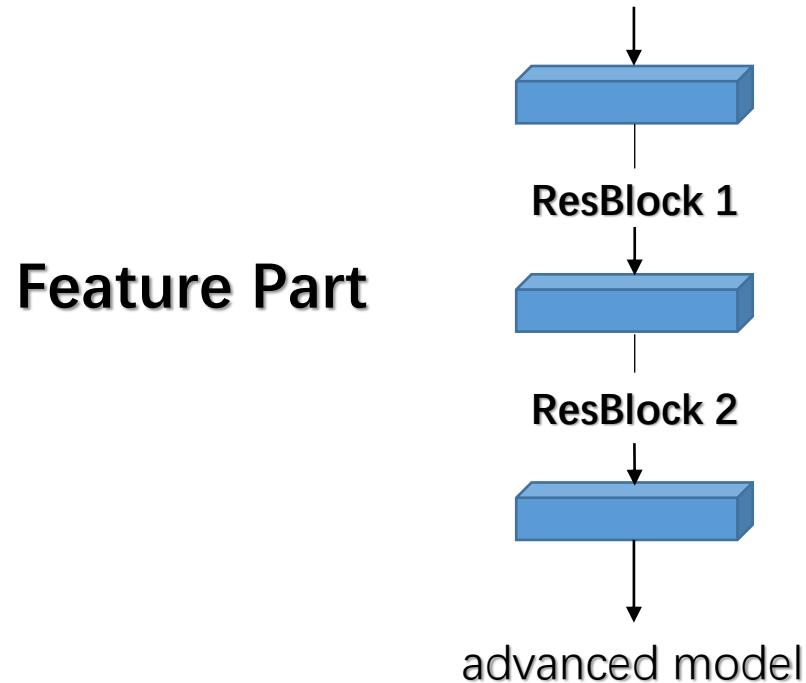
Bridge



I. Project I: Colorization

B. Updated Version

- UNet Structure [2015, Olaf Ronneberger]



I. Project I: Colorization

C. GAN based method

- **G-Net**

We could just use modified UNet

- **D-Net**

At current stage, D is simpler

```
self.conv0 = ConvBR(c_in=3,  
                    kernel_=  
                    weight_=  
self.conv1 = ConvBR(c_in=di,  
                    kernel_=  
                    weight_=  
self.conv2 = ConvBR(c_in=di,  
                    kernel_=  
                    weight_=  
self.conv3 = ResBasicBlock(  
  
self.conv4 = ResBasicBlock(  
  
self.conv5 = ConvBR(c_in=di,  
                    kernel_=  
                    weight_=
```

```
def forward(self, x):  
    # 256 -> 128  
    x = self.conv0(x)  
    # 128 -> 64  
    x = self.conv1(x)  
    # 64 -> 32  
    x = self.conv2(x)  
    # 32 -> 32  
    x = self.conv3(x)  
    # 32 -> 32  
    x = self.conv4(x)  
    # 32 -> 32  
    x = self.conv5(x)  
    return x
```

I. Project I: Colorization

C. GAN based method

- **G-Net**

We could just use modified UNet

- **D-Net**

At current stage, D is simpler

Question:

Why the output size is 32x32 after D?

Why it's not a single number

[no matter whether it's normal GAN or
WGAN]

we'll talk later

```
self.conv0 = ConvBR(c_in=3,  
                    kernel_=  
                    weight_=  
self.conv1 = ConvBR(c_in=di,  
                    kernel_=  
                    weight_=  
self.conv2 = ConvBR(c_in=di,  
                    kernel_=  
                    weight_=  
self.conv3 = ResBasicBlock(  
  
self.conv4 = ResBasicBlock(  
  
self.conv5 = ConvBR(c_in=di,  
                    kernel_=  
                    weight_=
```

```
def forward(self, x):  
    # 256 -> 128  
    x = self.conv0(x)  
    # 128 -> 64  
    x = self.conv1(x)  
    # 64 -> 32  
    x = self.conv2(x)  
    # 32 -> 32  
    x = self.conv3(x)  
    # 32 -> 32  
    x = self.conv4(x)  
    # 32 -> 32  
    x = self.conv5(x)  
    return x
```

I. Project I: Colorization

C. GAN based method

- Loss(es)

D-Net

WGAN-GP {
adversarial loss [WGAN]
gradient penalty [-GP]

G-Net

WGAN: adversarial loss

pixel level loss: L1 (texture sensitive) / L2 (color sensitive) [both not enough]

perceptual loss

tv loss

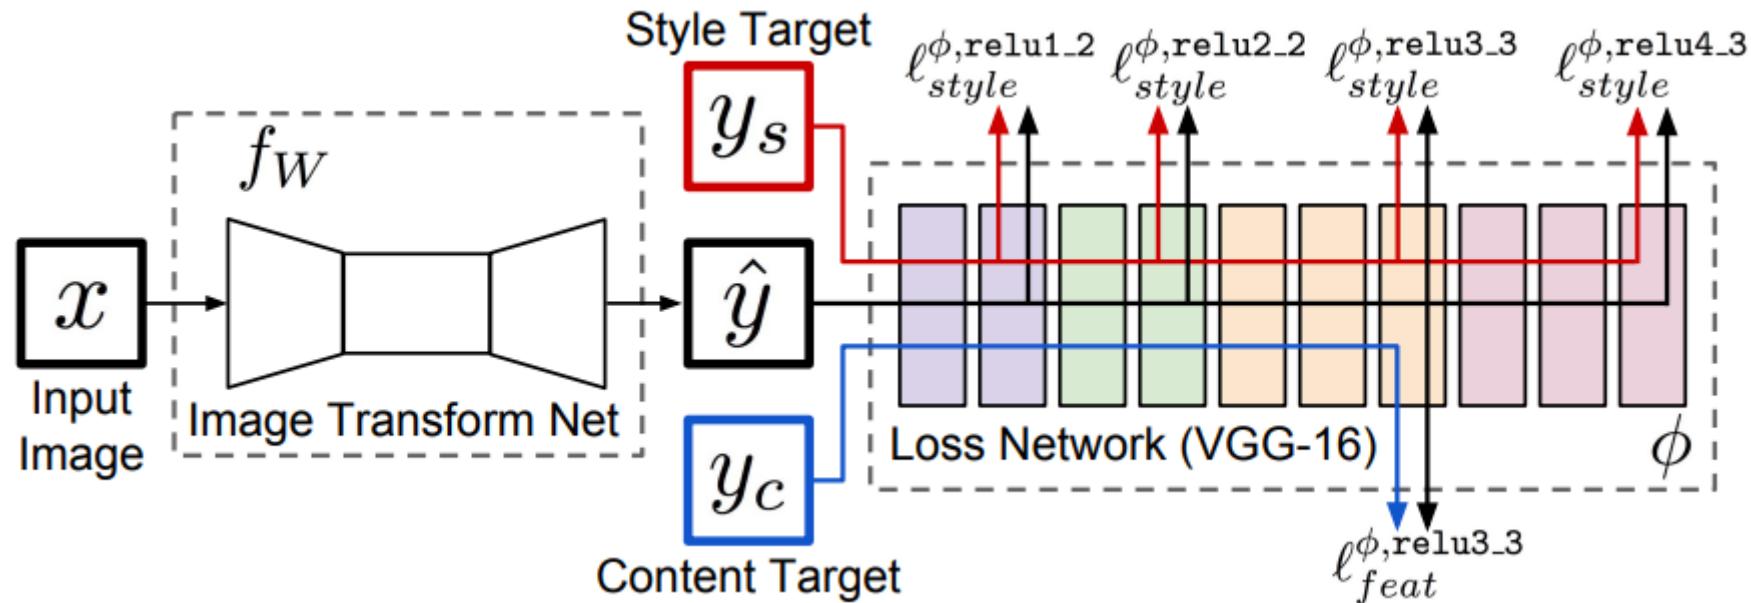
} which can also help to solve the problem of chessboard pattern

I. Project I: Colorization

D. Tricky Losses: Perceptual Loss

Perceptual Losses for Real-Time Style Transfer and Super-Resolution

[Justin Johnson, 2016, [link](#)]



I. Project I: Colorization

D. Perceptual Loss

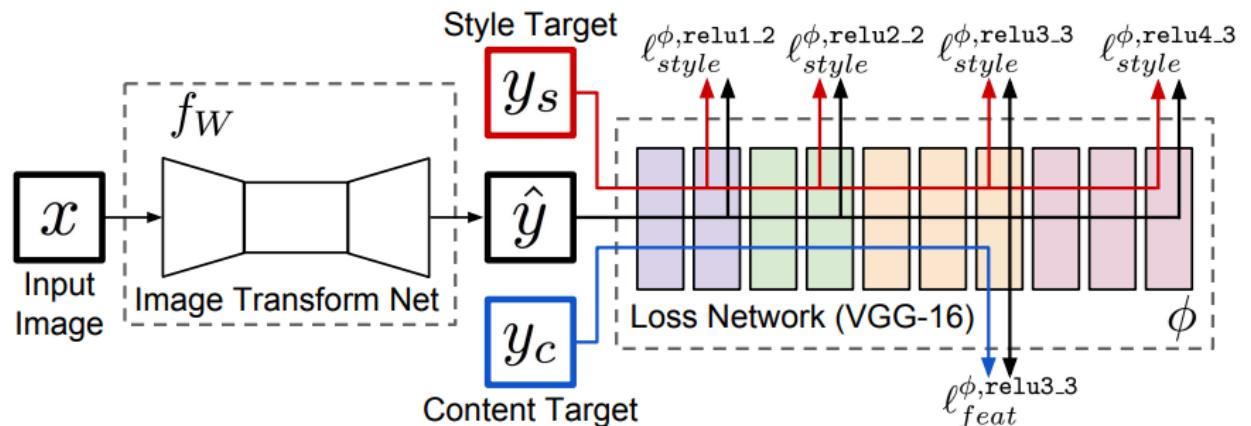
- Feature loss

$$\ell_{feat}^{\phi,j}(\hat{y}, y) = \frac{1}{C_j H_j W_j} \|\phi_j(\hat{y}) - \phi_j(y)\|_2^2$$

- Style reconstruction loss

$$G_j^\phi(x)_{c,c'} = \frac{1}{C_j H_j W_j} \sum_{h=1}^{H_j} \sum_{w=1}^{W_j} \phi_j(x)_{h,w,c} \phi_j(x)_{h,w,c'}$$

$$\ell_{style}^{\phi,j}(\hat{y}, y) = \|G_j^\phi(\hat{y}) - G_j^\phi(y)\|_F^2$$



Gram matrix

$$\mathbf{G} = \mathbf{X}^T \mathbf{X} = \begin{pmatrix} \mathbf{x}_1^T \\ \vdots \\ \mathbf{x}_m^T \end{pmatrix} (\mathbf{x}_1 \quad \cdots \quad \mathbf{x}_m), \quad G_{ij} = \mathbf{x}_i^T \mathbf{x}_j$$

I. Project I: Colorization

D. Perceptual Loss

- Feature loss

$$\ell_{feat}^{\phi,j}(\hat{y}, y) = \frac{1}{C_j H_j W_j} \|\phi_j(\hat{y}) - \phi_j(y)\|_2^2$$

- Style reconstruction loss

$$G_j^\phi(x)_{c,c'} = \frac{1}{C_j H_j W_j} \sum_{h=1}^{H_j} \sum_{w=1}^{W_j} \phi_j(x)_{h,w,c} \phi_j(x)_{h,w,c'}$$

$$\ell_{style}^{\phi,j}(\hat{y}, y) = \|G_j^\phi(\hat{y}) - G_j^\phi(y)\|_F^2$$

```
# calculate gram matrices for style feature layer maps we care about
style_features = vgg(style)
style_gram = [utils.gram(fmap) for fmap in style_features]
```

how to get gram

```
# calculate style loss
y_hat_gram = [utils.gram(fmap) for fmap in y_hat_features]
style_loss = 0.0
for j in range(4):
    style_loss += loss_mse(y_hat_gram[j], style_gram[j])
```

I. Project I: Colorization

D. Perceptual Loss

- Feature loss

$$\ell_{feat}^{\phi,j}(\hat{y}, y) = \frac{1}{C_j H_j W_j} \|\phi_j(\hat{y}) - \phi_j(y)\|_2^2$$

- Style reconstruction loss

$$G_j^\phi(x)_{c,c'} = \frac{1}{C_j H_j W_j} \sum_{h=1}^{H_j} \sum_{w=1}^{W_j} \phi_j(x)_{h,w,c} \phi_j(x)_{h,w,c'}$$

$$\ell_{style}^{\phi,j}(\hat{y}, y) = \|G_j^\phi(\hat{y}) - G_j^\phi(y)\|_F^2$$

```
torch.bmm(input, mat2, *, out=None) → Tensor
```

Performs a batch matrix-matrix product of matrices stored in `input` and `mat2`.

`input` and `mat2` must be 3-D tensors each containing the same number of matrices.

If `input` is a $(b \times n \times m)$ tensor, `mat2` is a $(b \times m \times p)$ tensor, `out` will be a $(b \times n \times p)$ tensor

$$\text{out}_i = \text{input}_i @ \text{mat2}_i$$

```
# Calculate Gram matrix (G = FF^T)
def gram(x):
    (bs, ch, h, w) = x.size()
    f = x.view(bs, ch, w*h)
    f_T = f.transpose(1, 2)
    G = f.bmm(f_T) / (ch * h * w)
```

I. Project I: Colorization

D. Perceptual Loss

- Feature loss

$$\ell_{feat}^{\phi,j}(\hat{y}, y) = \frac{1}{C_j H_j W_j} \|\phi_j(\hat{y}) - \phi_j(y)\|_2^2$$

Much more popular because
style reconstruction loss is much heavier

- Style reconstruction loss

$$G_j^\phi(x)_{c,c'} = \frac{1}{C_j H_j W_j} \sum_{h=1}^{H_j} \sum_{w=1}^{W_j} \phi_j(x)_{h,w,c} \phi_j(x)_{h,w,c'}$$

$$\ell_{style}^{\phi,j}(\hat{y}, y) = \|G_j^\phi(\hat{y}) - G_j^\phi(y)\|_F^2$$

I. Project I: Colorization

D. Tricky Loss: tv Loss

- **total variation loss/regularizer**
[image gradient loss/regularizer]

Theory:

$$\mathcal{R}_{V^\beta}(f) = \int_{\Omega} \left(\left(\frac{\partial f}{\partial u}(u, v) \right)^2 + \left(\frac{\partial f}{\partial v}(u, v) \right)^2 \right)^{\frac{\beta}{2}} du dv$$

$$\mathcal{R}_{V^\beta}(\mathbf{x}) = \sum_{i,j} \left((x_{i,j+1} - x_{ij})^2 + (x_{i+1,j} - x_{ij})^2 \right)^{\frac{\beta}{2}}$$

$\beta = 1$

Implementation: practical way

```
tv_loss = torch.mean(torch.abs(pred[:, :, :-1, :] -  
                           pred[:, :, 1:, :])) + \  
                  torch.mean(torch.abs(pred[:, :, :, :-1] -  
                           pred[:, :, :, 1:])))
```

II. Project II: Face Frontalization



II. Project II: Face Frontalization

- Paper: Two-Pathway GAN [2017] (TP-GAN)

Authors: Rui Huang etc.

Link: [address](#)

Other resources: [DR-GAN](#)

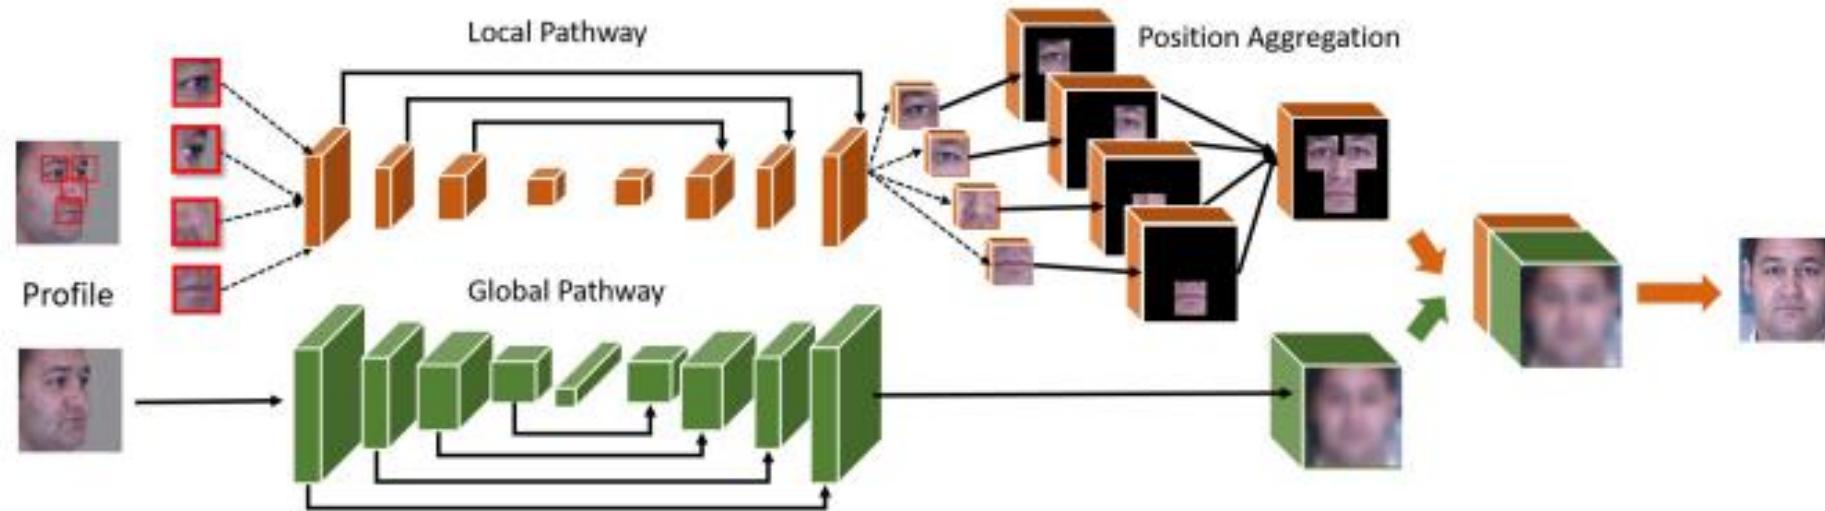


I. Project II: Face Frontalization

E. Modified Backbone

- **Several Parts**

Two-pathway Generator Network

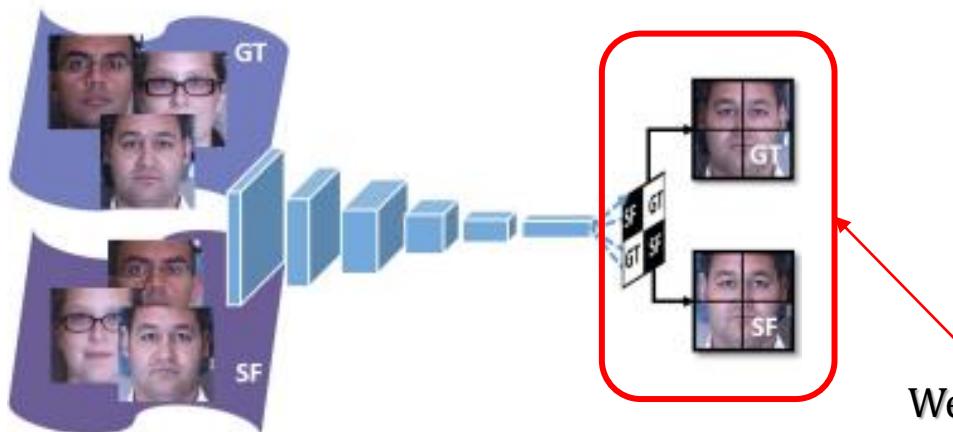


I. Project II: Face Frontalization

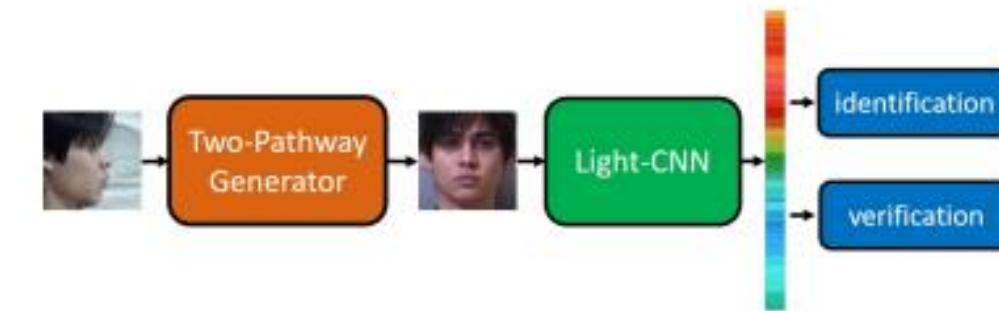
E. Modified Backbone

- **Several Parts**

Discriminator Network



Recognition via Generation



We meet this scene again
[This is called PatchGAN,
we'll cover this next time]

I. Project II: Face Frontalization

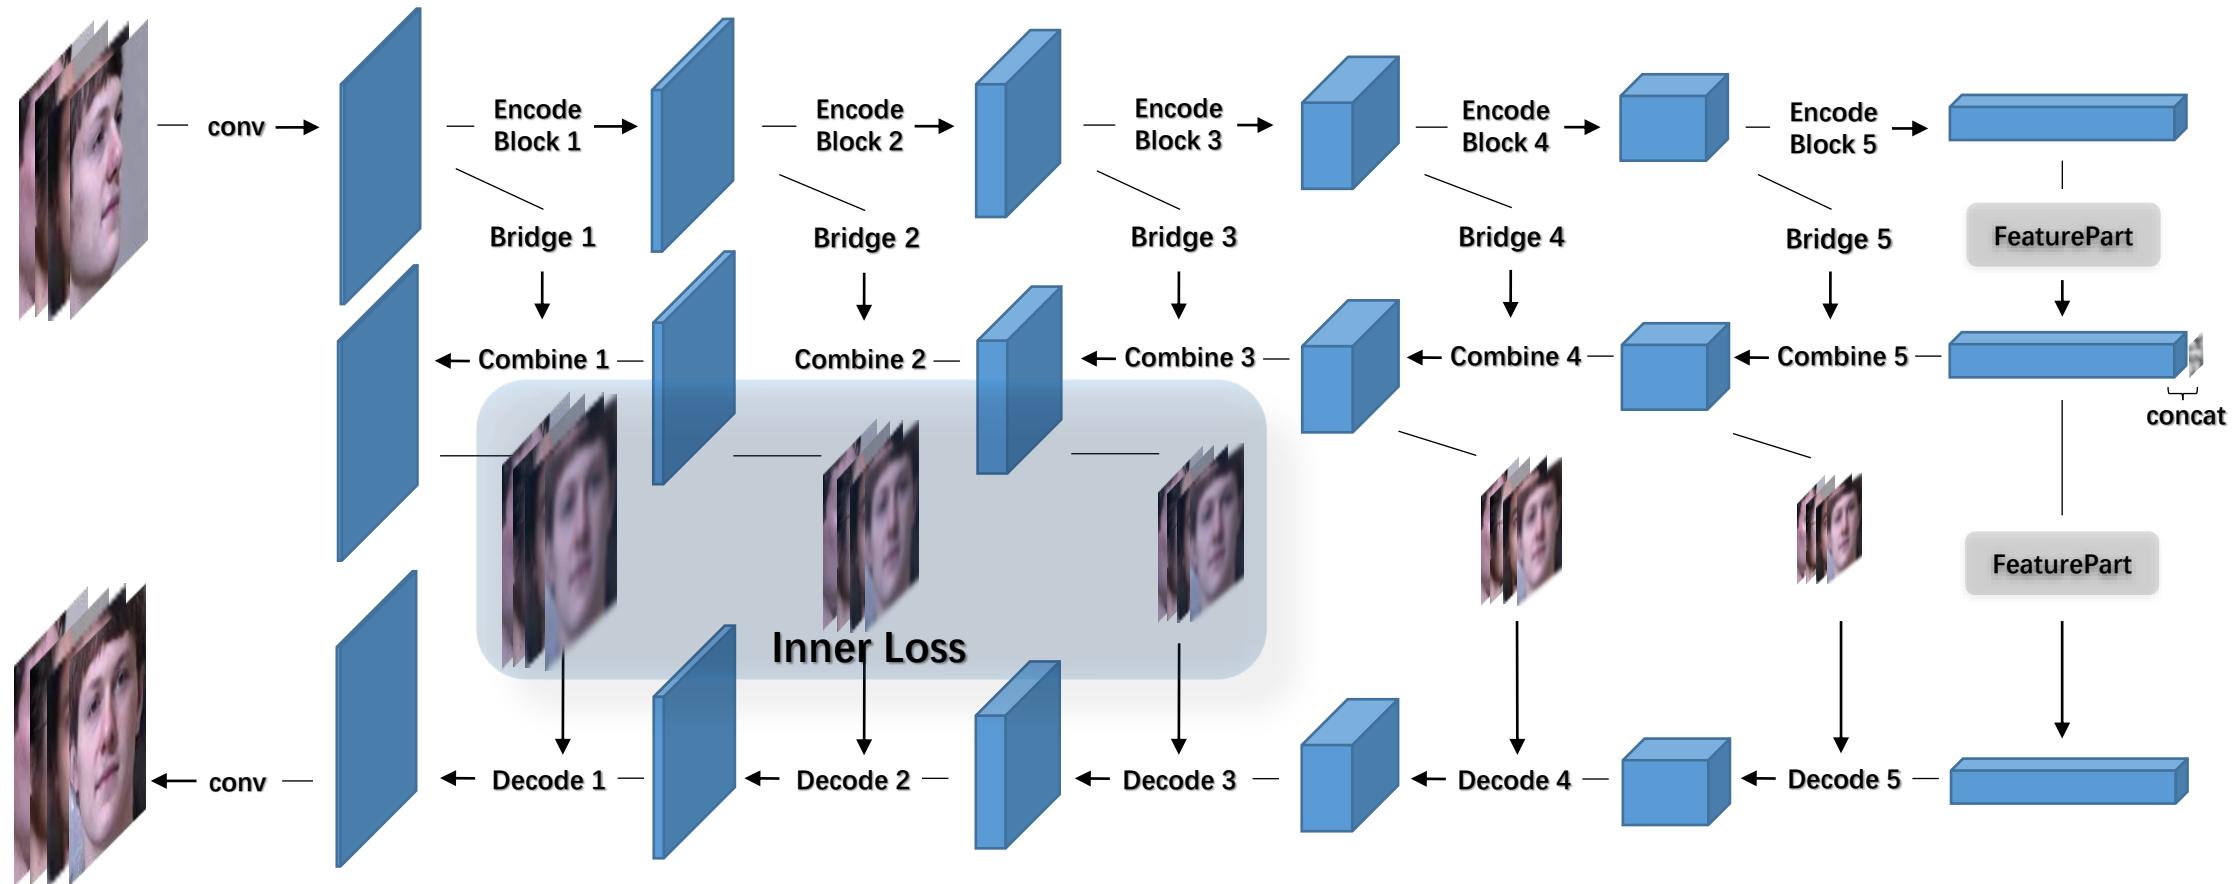
E. Modified Backbone

- **Key Parts: Backbone of Global Pathway: Updated from UNet**

I. Project II: Face Frontalization

E. Modified Backbone

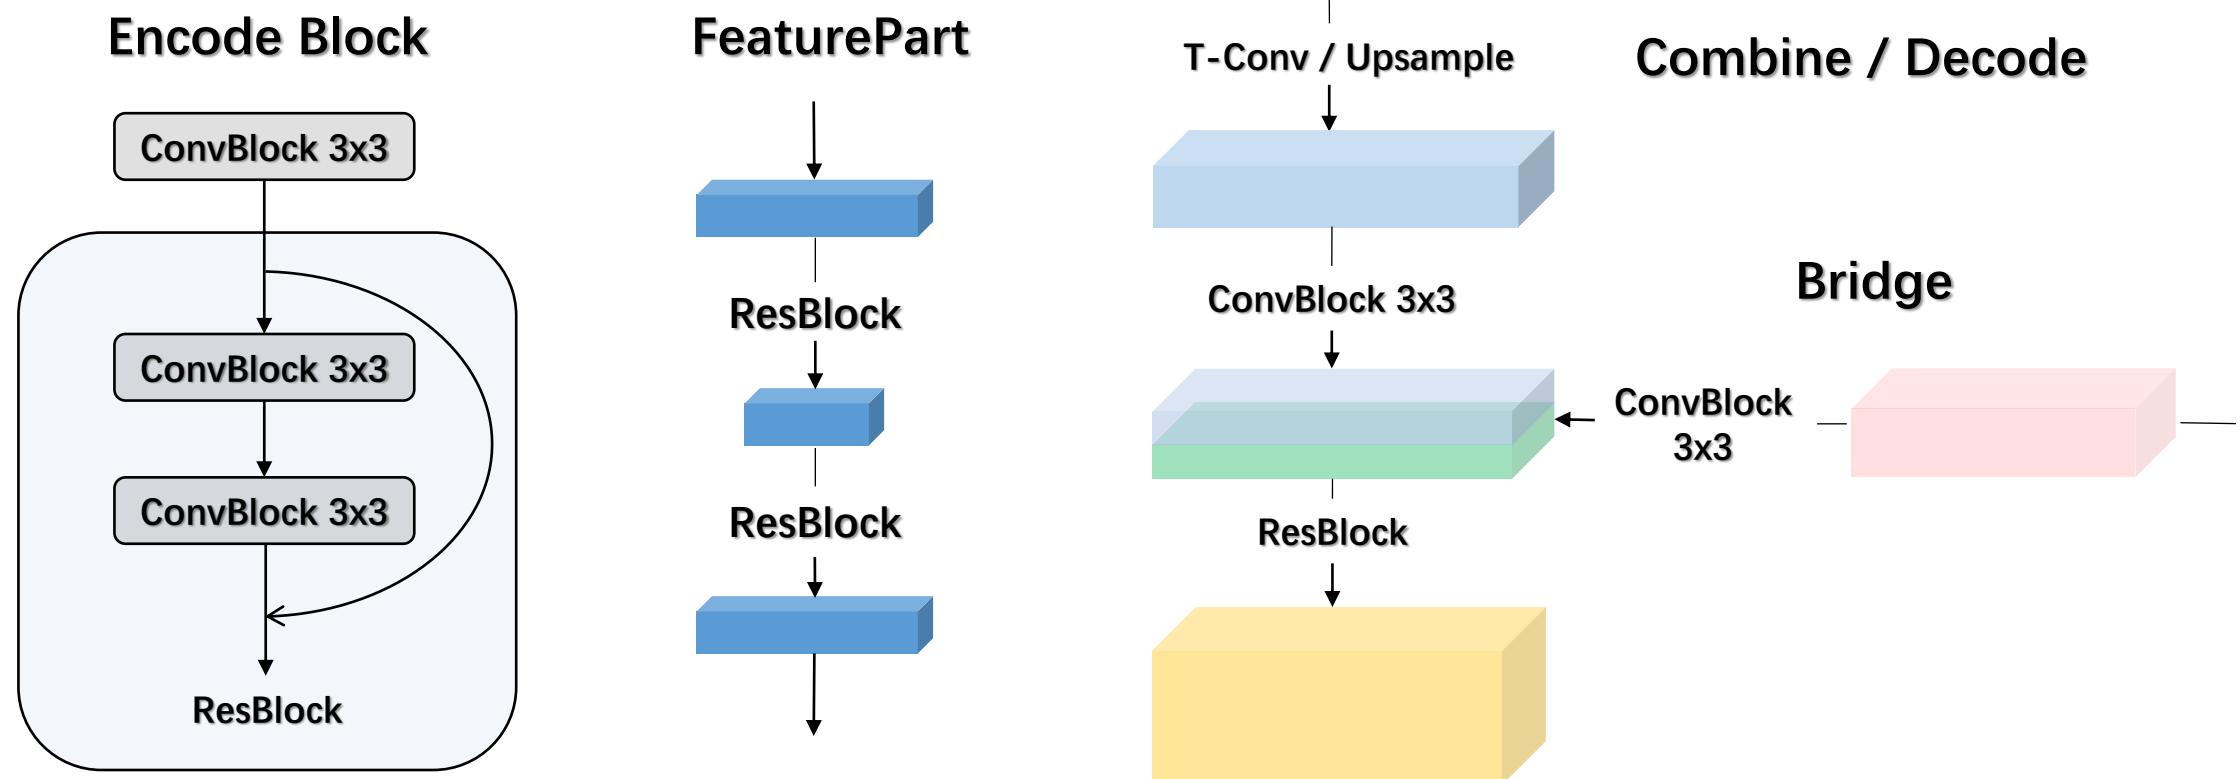
- **Key Parts: Backbone of Global Pathway: Updated from UNet**



I. Project II: Face Frontalization

E. Modified Backbone

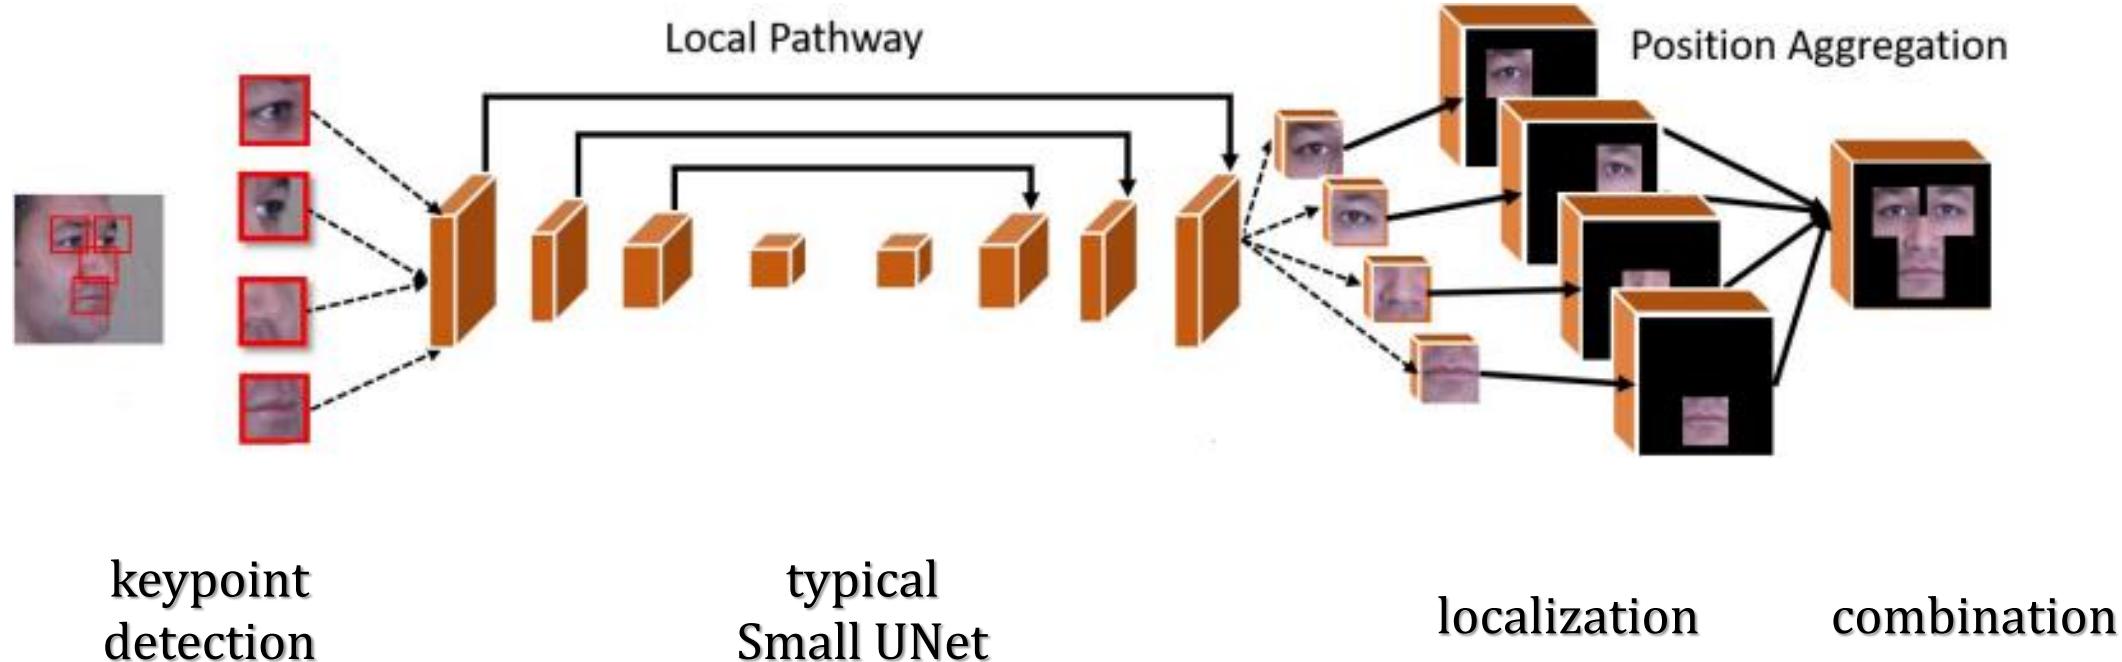
- **Key Parts: Details of Global Pathway**



I. Project II: Face Frontalization

E. Modified Backbone

- **Key Parts: Local Pathway**

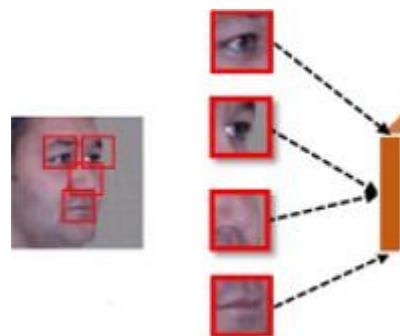


I. Project II: Face Frontalization

E. Modified Backbone

- **Key Parts: Local Pathway**

[not our target in this class]



Method 1: [Dlib](#) (this is official website in C++)

[[installation](#), install dlib before use]

[e.g., a good example to learn how to use in Python]

2: MTCNN ([paper](#), [code](#)) [old fashioned in TF])

[e.g., a good example we can directly use in PyTorch]

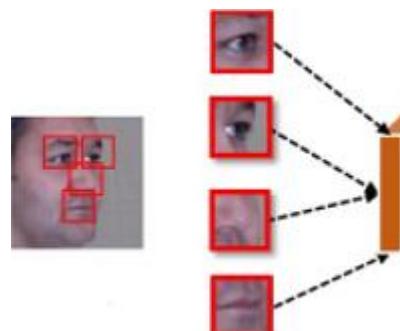
Details (original): [[any ideas to upgrade?](#)]

1. 5 landmarks: 2 eye-centers / 1 nose / 2 corners of the mouth
2. Crop a little patch around each detected landmark
by default: img size: 128 / eye: 40 / nose: 40-32 / mouth: 48-32
3. Each path put into local pathway UNet

I. Project II: Face Frontalization

E. Modified Backbone

- **Key Parts: Local Pathway**



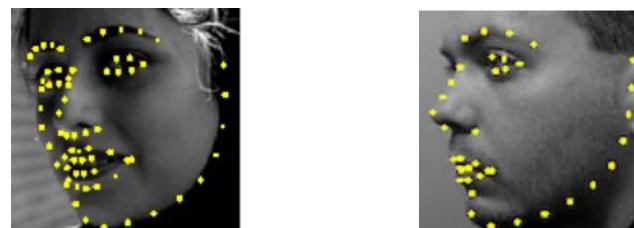
keypoint
detection

Details (original): [any ideas to upgrade?]

1. 5 landmarks: 2 eye-centers / 1 nose / 2 corners of the mouth
2. Crop a little patch around each detected landmark
by default: img size: 128 / eye: 40 / nose: 40-32 / mouth: 48-32
3. Each path put into local pathway UNet

Details:

1. **more** landmarks
2. Crop an **accurate** patch around each detected landmark

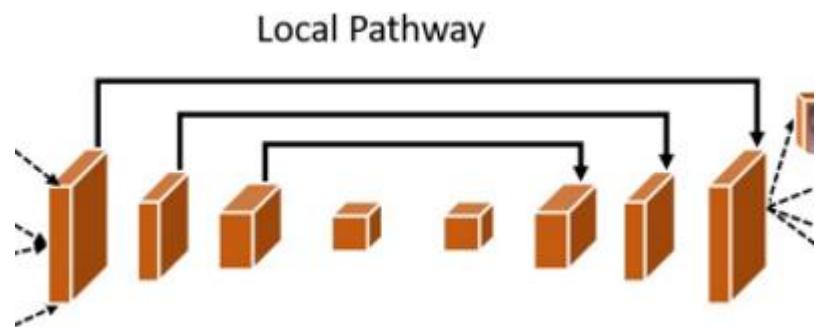


3. Each path put into local pathway UNet

I. Project II: Face Frontalization

E. Modified Backbone

- **Key Parts: Local Pathway**



typical
Small UNet

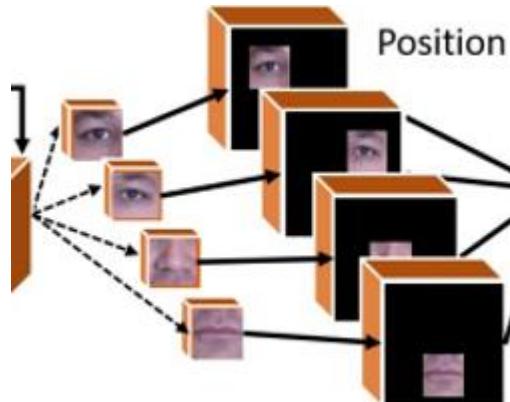
Details:

For each patch, we instantiate a network
5 patches → 5 instantiated network

I. Project II: Face Frontalization

E. Modified Backbone

- **Key Parts: Local Pathway**



Details

For each frontalized little patch, padding to greater tensors with same size

```
def forward(self, f_left_eye, f_right_eye, f_nose, f_mouth):
    f_left_eye = torch.nn.functional.pad(f_left_eye,
                                         (leye_w - EYE_W // 2 - 1,
                                         input_img_size - (leye_w + EYE_W // 2 - 1),
                                         leye_h - EYE_H // 2 - 1,
                                         input_img_size - (leye_h + EYE_H // 2 - 1))) # input
    # left
    # right
    # top
    # bottom

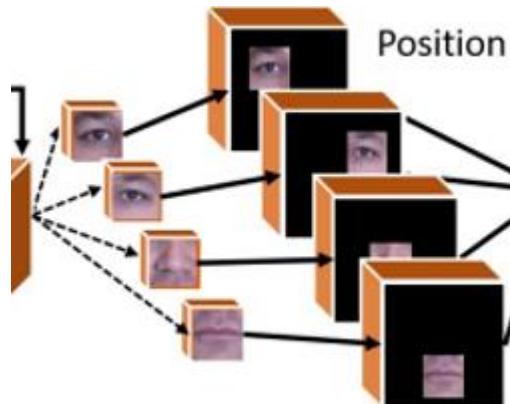
    f_right_eye = torch.nn.functional.pad(f_right_eye,
                                         (reye_w - EYE_W // 2 - 1,
                                         input_img_size - (reye_w + EYE_W // 2 - 1),
                                         reye_h - EYE_H // 2 - 1,
```

How to localize from where to pad?

I. Project II: Face Frontalization

E. Modified Backbone

- **Key Parts: Local Pathway**



localization

Details

For each frontalized little patch, padding to greater tensors with same size

```
def forward(self, f_left_eye, f_right_eye, f_nose, f_mouth):
    f_left_eye = torch.nn.functional.pad(f_left_eye,
                                         (leye_w - EYE_W // 2 - 1,
                                         input_img_size - (leye_w + EYE_W // 2 - 1),
                                         leye_h - EYE_H // 2 - 1,
                                         input_img_size - (leye_h + EYE_H // 2 - 1))) # input
                                         # left
                                         # right
                                         # top
                                         # bottom

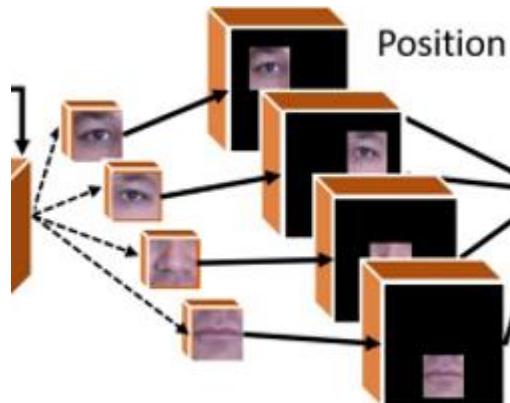
    f_right_eye = torch.nn.functional.pad(f_right_eye,
                                         (reye_w - EYE_W // 2 - 1,
                                         input_img_size - (reye_w + EYE_W // 2 - 1),
                                         reye_h - EYE_H // 2 - 1,
```

How to localize from where to pad?

I. Project II: Face Frontalization

E. Modified Backbone

- **Key Parts: Local Pathway**



Details

For each frontalized little patch, padding to greater tensors with same size

```
def forward(self, f_left_eye, f_right_eye, f_nose, f_mouth):
    f_left_eye = torch.nn.functional.pad(f_left_eye,
                                         (leye_w - EYE_W // 2 - 1,
                                         input_img_size - (leye_w + EYE_W // 2 - 1),
                                         leye_h - EYE_H // 2 - 1,
                                         input_img_size - (leye_h + EYE_H // 2 - 1))) # input
                                         # left
                                         # right
                                         # top
                                         # bottom

    f_right_eye = torch.nn.functional.pad(f_right_eye,
                                         (reye_w - EYE_W // 2 - 1,
                                         input_img_size - (reye_w + EYE_W // 2 - 1),
                                         reye_h - EYE_H // 2 - 1,
```

localization

```
_parser.add_argument("--input_img_size", default=128, type=int)
_parser.add_argument("--EYE_W", default=40, type=int)
_parser.add_argument("--EYE_H", default=40, type=int)
_parser.add_argument("--NOSE_W", default=40, type=int)
_parser.add_argument("--NOSE_H", default=32, type=int)
_parser.add_argument("--reye_w", default=44, type=int)
_parser.add_argument("--reye_h", default=41, type=int)
_parser.add_argument("--leye_w", default=86, type=int)
_parser.add_argument("--leye_h", default=41, type=int)
```

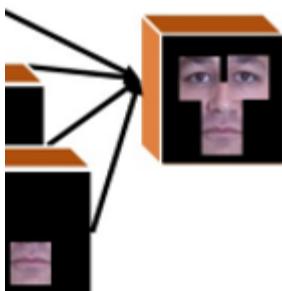
I. Project II: Face Frontalization

E. Modified Backbone

- **Key Parts: Local Pathway**

Details

Combine/Merge padded frontalized patch/tensor as a whole



combination

```
return torch.max(  
    torch.stack([f_left_eye, f_right_eye, f_nose, f_mouth], dim=0),  
    dim=0  
) [0]
```

1. `torch.stack([a, b, c, d], dim=0)`, 在第dim维入栈

So if `A` and `B` are of shape (3, 4):

- `torch.cat([A, B], dim=0)` will be of shape (6, 4)
- `torch.stack([A, B], dim=0)` will be of shape (2, 3, 4)

Left eye: [1, 64, 128, 128] Left mouth: [1, 64, 128, 128]

[] Stacked tensor: [5, 1, 64, 128, 128]

I. Project II: Face Frontalization

E. Modified Backbone

- **Key Parts: Local Pathway**

Details

Combine/Merge padded frontalized patch/tensor as a whole



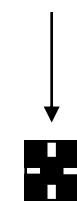
combination

```
return torch.max(  
    torch.stack([f_left_eye, f_right_eye, f_nose, f_mouth], dim=0),  
    dim=0  
) [0]
```

2. `torch.max(tensor, dim=0)`, 沿第dim找最大值并存出来

```
>>> a = torch.randn(2, 3, 4)  
>>> a  
tensor([[[ -1.3002,   0.0548,   1.0300,   0.6392],  
        [-0.1983,   0.1366,  -0.0305,  -0.0889],  
        [ 2.6973,   1.4979,  -0.2619,  -0.3422]],  
  
        [[ -0.9631,   0.2432,  -0.0634,   0.2243],  
        [-0.6535,   0.5095,   1.6183,   0.3808],  
        [ -1.0251,   2.0204,   2.4084,   1.6377]]])  
>>> torch.max(a, 0)  
torch.return_types.max(  
values=tensor([[-0.9631,  0.2432,  1.0300,  0.6392],  
            [-0.1983,  0.5095,  1.6183,  0.3808],  
            [ 2.6973,  2.0204,  2.4084,  1.6377]]),
```

`max([`  `,`  `,`  `,`  `], dim=0)`



I. Project II: Face Frontalization

E. Modified Backbone

- **Key Parts: Local Pathway**

Details

Combine/Merge padded frontalized patch/tensor as a whole



combination

```
return torch.max(  
    torch.stack([f_left_eye, f_right_eye, f_nose, f_mouth], dim=0),  
    dim=0  
) [0]
```

3. `torch.max(tensor, dim=0)[0]`, `torch.max`返回2个值，最大值[0] 和 对应位置[1]

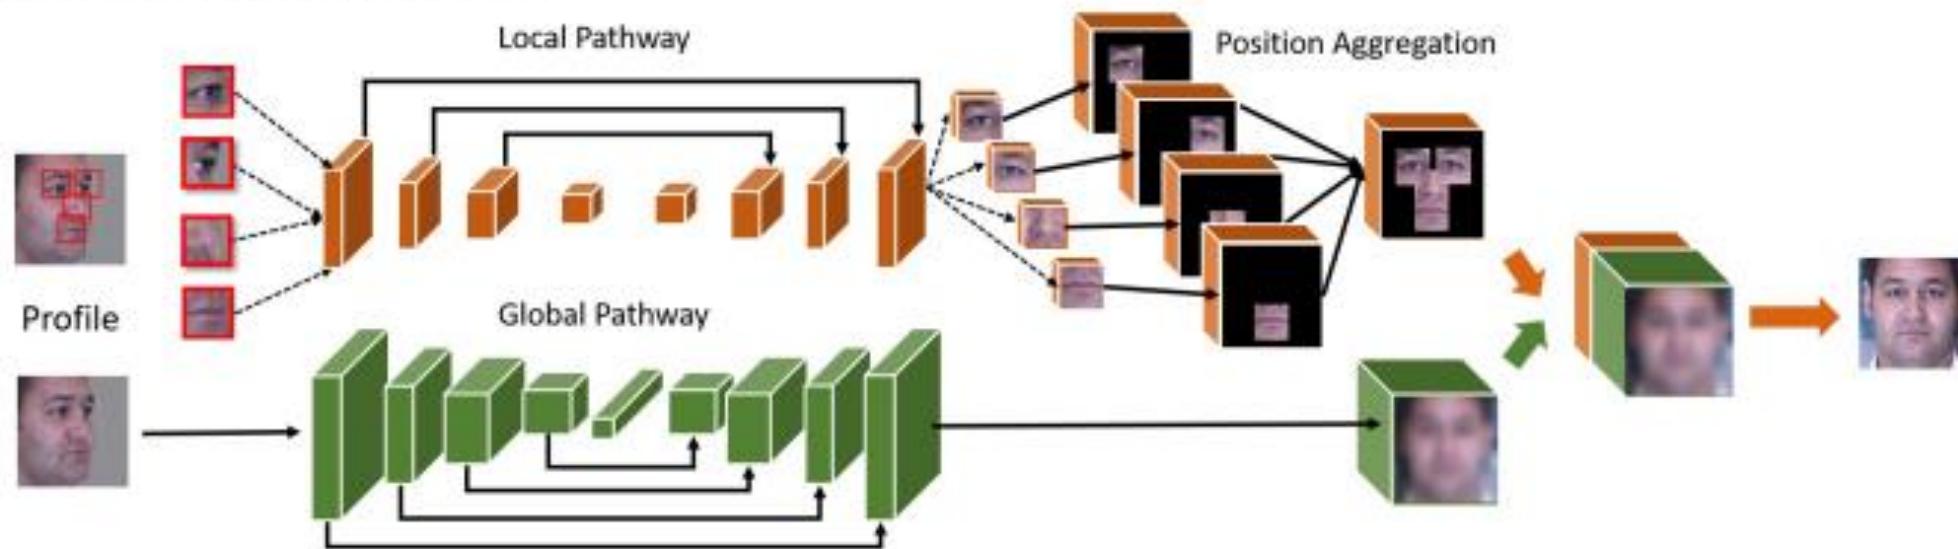
```
>>> a = torch.randn(2, 3, 4)  
>>> a  
tensor([[[ -1.3002,   0.0548,   1.0300,   0.6392],  
        [-0.1983,   0.1366,  -0.0305,  -0.0889],  
        [ 2.6973,   1.4979,  -0.2619,  -0.3422]],  
        [[ -0.9631,   0.2432,  -0.0634,   0.2243],  
        [-0.6535,   0.5095,   1.6183,   0.3808],  
        [ -1.0251,   2.0204,   2.4084,   1.6377]]])  
>>> torch.max(a, 0)  
torch.return_types.max(  
values=tensor([-0.9631,  0.2432,  1.0300,  0.6392],  
[-0.1983,  0.5095,  1.6183,  0.3808],  
[ 2.6973,  2.0204,  2.4084,  1.6377])),  
indices=tensor([[1, 1, 0, 0],  
[0, 1, 1, 1],  
[0, 1, 1, 1]]))
```

I. Project II: Face Frontalization

E. Modified Backbone

- **Key Parts: Global Pathway: UNet Upgrade + Combination with Local**

Two-pathway Generator Network



I. Project II: Face Frontalization

F. Tricky Losses: Symmetrical Loss + Extra Loss

- **Symmetrical Loss**

Definition:

Common faces are symmetrical at most time.

Why not use this feature as a training target

How to use:

```
class SymLoss(nn.modules.loss._Loss):

    def __init__(self):
        super(SymLoss, self).__init__()
        self.l1 = nn.L1Loss(reduction="sum")

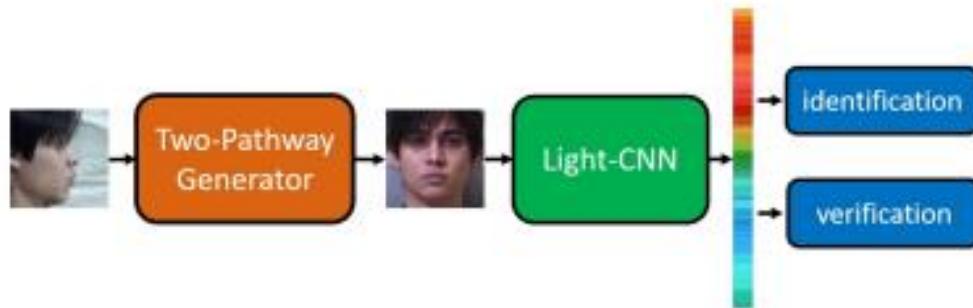
    def forward(self, x):
        half_w = x.size(-1) // 2
        return self.l1(x[..., :half_w], x[..., half_w:].flip(-1))
```

I. Project II: Face Frontalization

F. Tricky Losses: Symmetrical Loss + Extra Loss

- **Extra Loss: Recognition Loss**

Recognition via Generation



Reason:

We hope the generated frontalized faces are faces of the same people

How to use:

1. Push generated frontalized faces into face recognition network to get corresponding features
2. Then compare the features as Perceptual Loss

I. Project II: Face Frontalization

F. Tricky Losses

- **Summary**

You can use as many losses as possible.

The final loss is the combination of different partial losses which are your learning targets

You need to find out weights corresponding to each loss in order their values are balanced

III. Project I: Revisit



III. Project I: Revisit

G. Temporal Consistency



III. Project I: Revisit

G. Temporal Consistency



III. Project I: Revisit

G. Temporal Consistency

- **How to constrain the temporal consistency**

Optical Flow: Definition

the pattern of apparent motion of image objects between two consecutive frames caused by the movement of object or camera.

2D vector field where each vector is a displacement vector showing the movement of points from first frame to second



III. Project I: Revisit

G. Temporal Consistency

- **How to constrain the temporal consistency**

Optical Flow: Assumption

1. The pixel intensities of an object do not change between consecutive frames.



III. Project I: Revisit

G. Temporal Consistency

- How to constrain the temporal consistency

Optical Flow: Calculation

$$I(x, y, t) = I(x + dx, y + dy, t + dt)$$

$$= I(x, y, t) + \frac{\partial I}{\partial x} dx + \frac{\partial I}{\partial y} dy + \frac{\partial I}{\partial t} dt + o^n$$

⇒

$$\frac{\partial I}{\partial x} \frac{dx}{dt} + \frac{\partial I}{\partial y} \frac{dy}{dt} + \frac{\partial I}{\partial t} \frac{dt}{dt} = \frac{\partial I}{\partial x} \frac{dx}{dt} + \frac{\partial I}{\partial y} \frac{dy}{dt} + \frac{\partial I}{\partial t} = 0$$

⇒

$I_x u + I_y v + I_t = 0$ where (u, v) is the optical vector and
 I_x, I_y and I_t can get from two consecutive frames

Difficulty:

1 equation with 2 parameters (u, v)

So, it means we still need at least another equation

III. Project I: Revisit

G. Temporal Consistency

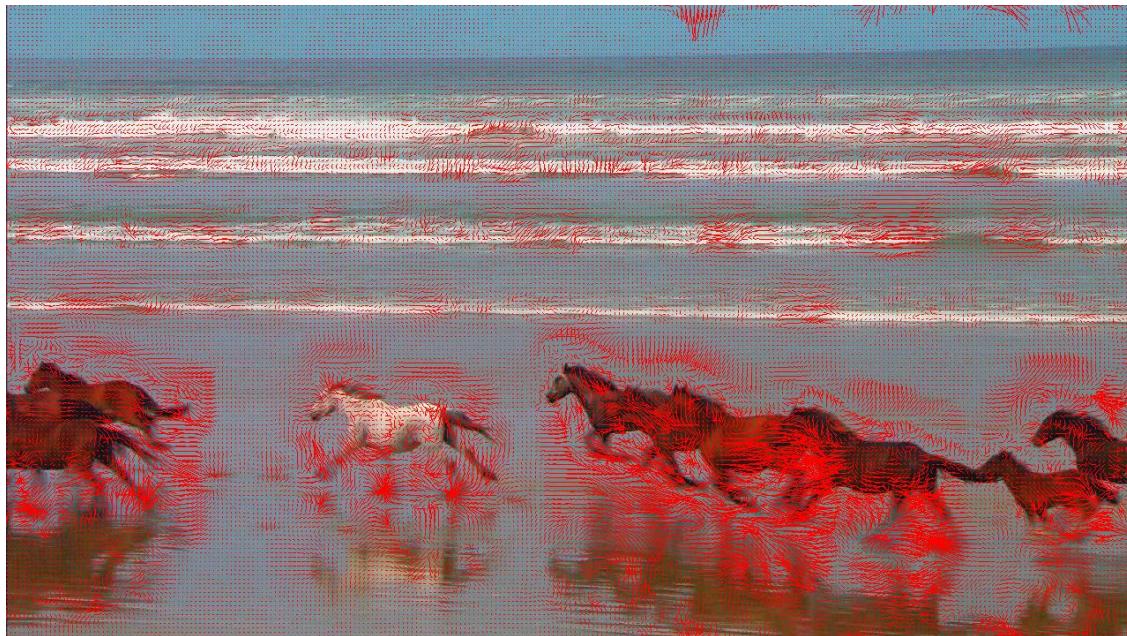
- **How to constrain the temporal consistency**

Optical Flow: Assumption

1. The pixel intensities of an object do not change between consecutive frames.
2. Neighboring pixels have similar motion

(There are many methods to get other constraints to find out (u, v) , here we only provide the most common method.)

You need to understand it before your interview



III. Project I: Revisit

G. Temporal Consistency

- **How to constrain the temporal consistency**

Optical Flow: Lucas-Kanade (LK) Algorithm [1981]

Critical idea:

Pixels inside a $N \times N$ window (neighboring pixels) have the same movement

$$\mathbf{A}^T \mathbf{v} = \mathbf{b}$$

$$\begin{aligned} I_x(q_1)u + I_y(q_1)v &= -I_t(q_1) \\ I_x(q_2)u + I_y(q_2)v &= -I_t(q_2) \\ &\vdots \\ I_x(q_n)u + I_y(q_n)v &= -I_t(q_n) \end{aligned} \Rightarrow \begin{aligned} \mathbf{A} &= \begin{bmatrix} I_x(q_1) & I_x(q_2) & \dots & I_x(q_n) \\ I_y(q_1) & I_y(q_2) & & I_y(q_n) \end{bmatrix} \\ \mathbf{v} &= \begin{bmatrix} u \\ v \end{bmatrix} \end{aligned}$$

$$\mathbf{b} = [-I_t(q_1) \quad -I_t(q_1) \quad \dots \quad -I_t(q_1)]^T$$

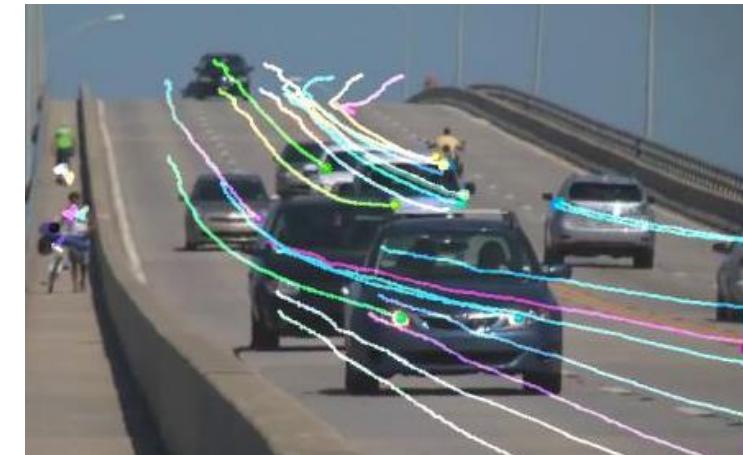
III. Project I: Revisit

G. Temporal Consistency

- **How to constrain the temporal consistency**

Optical Flow: Lucas-Kanade (LK) Algorithm [\[1981\]](#) (Sparse Optical Flow)

$$\begin{aligned}\boldsymbol{v} &= \begin{bmatrix} u \\ v \end{bmatrix} \\ &= \begin{bmatrix} \sum_i I_x(q_i)^2 & \sum_i I_x(q_i)I_y(q_i) \\ \sum_i I_x(q_i)I_y(q_i) & \sum_i I_y(q_i)^2 \end{bmatrix}^{-1} \begin{bmatrix} -\sum_i I_x(q_i)I_t(q_i) \\ -\sum_i I_y(q_i)I_t(q_i) \end{bmatrix}\end{aligned}$$



III. Project I: Revisit

G. Temporal Consistency

- **How to constrain the temporal consistency**

Optical Flow: Some other materials

Dense Optical Flow: Farneback Optical Flow (2003): (we can use this in light project)



Optical Flow in Deep Learning: (we can use this in real project)

Flownet 2.0: Evolution of Optical Flow Estimation with Deep Networks (Eddy Ilg, 2016)

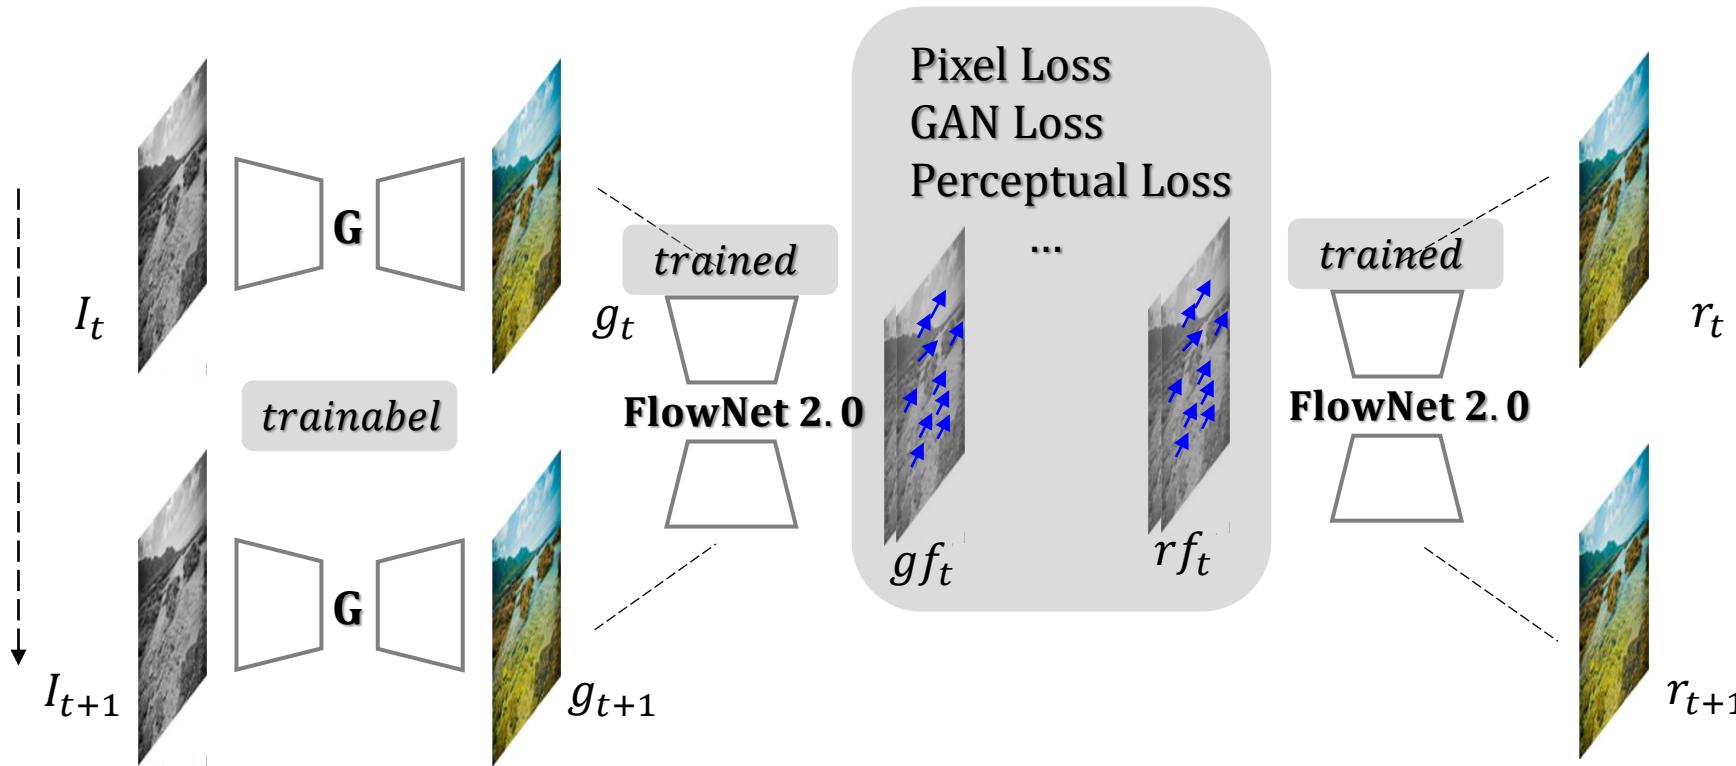
Related papers: [link](#)

III. Project I: Revisit

G. Temporal Consistency

- How to integrate optical flow into colorizing video

method 1: directly add optical flow constraint



III. Project I: Revisit

G. Temporal Consistency

- How to integrate optical flow into colorizing video

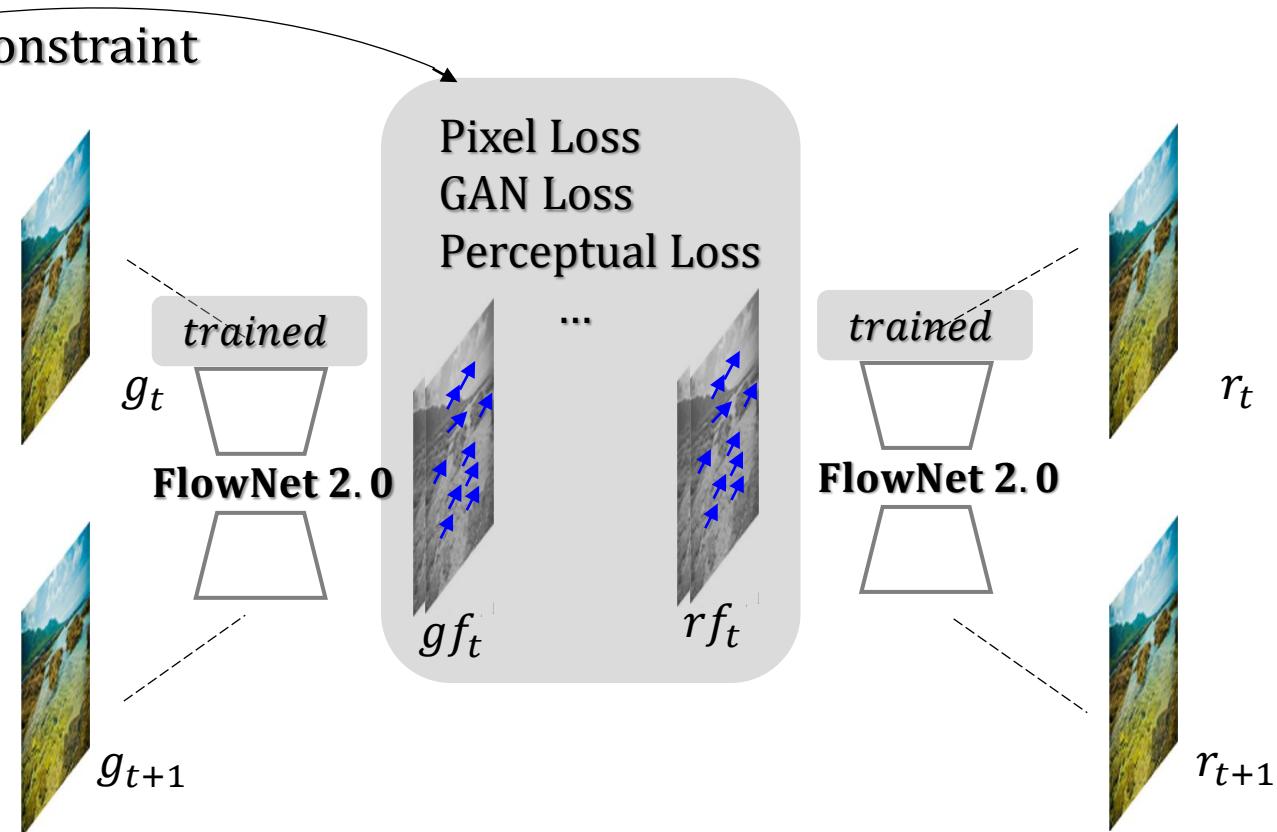
method 1: directly add optical flow constraint

Generalized tricks:

1. Optical flow result has 2 channels:
[displacement_x, displacement_y]
2. For **flow_loss**, we can make it more robust by warping the output like:

$$\mathcal{L}_{gf} = \text{losses_we_have_known} + \| gf_t(\mathbf{r}_t) - \mathbf{r}_{t+1} \|$$

An [example](#) to show how to implement
(from line 108-115)



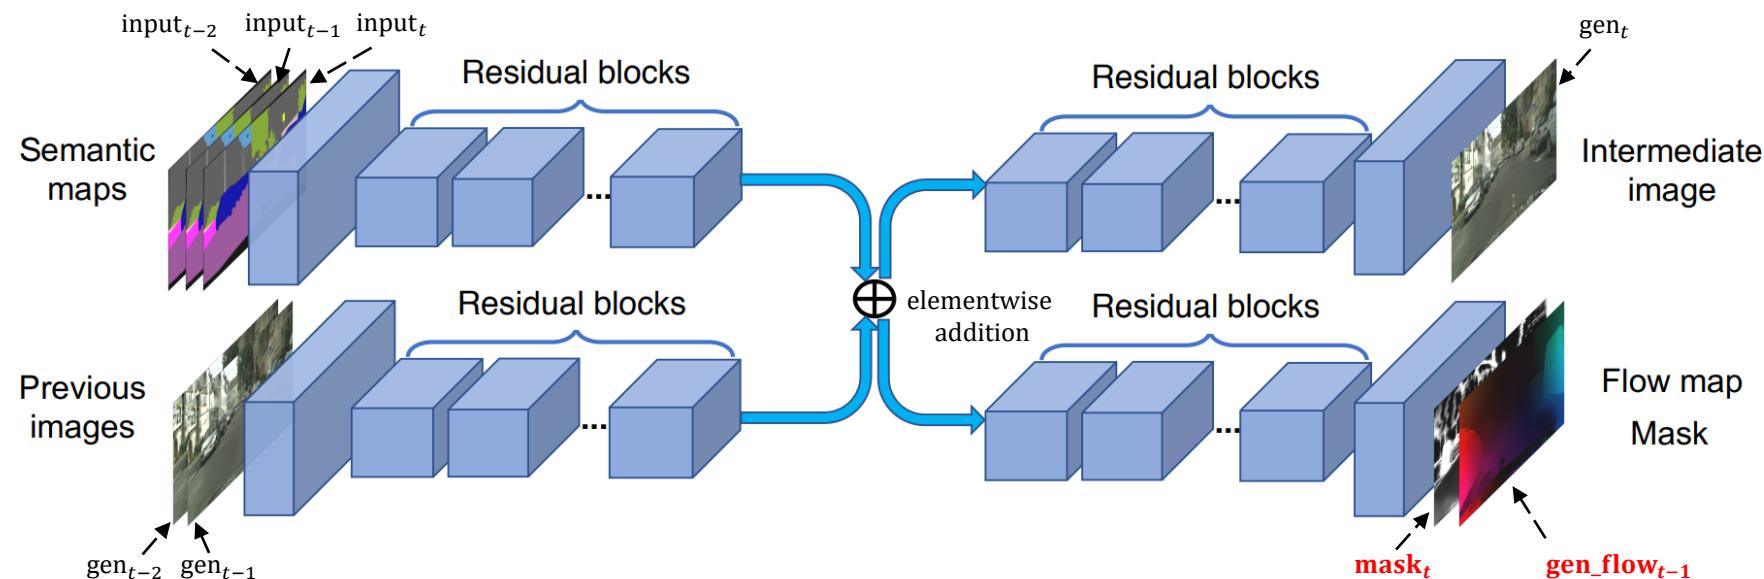
III. Project I: Revisit

G. Temporal Consistency

- **How to integrate optical flow into colorizing video**

method 2: use output as input

[vid2vid: Video-to-Video Synthesis (2018, Wang)]



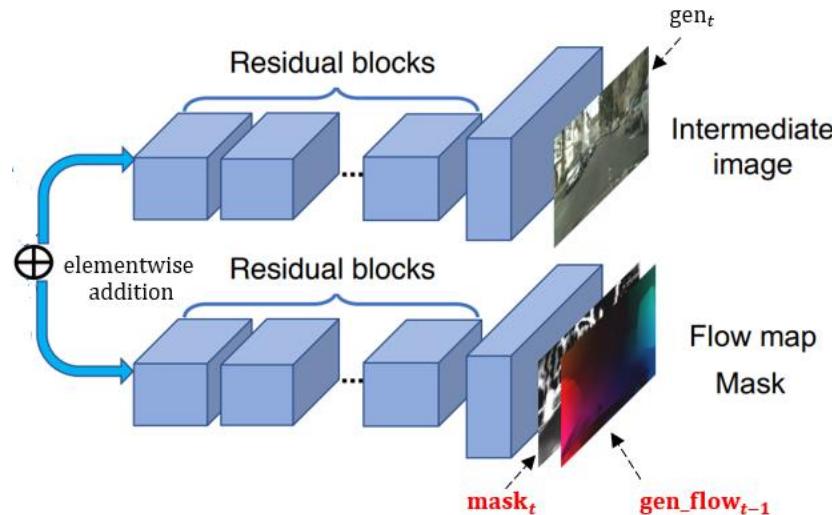
III. Project I: Revisit

G. Temporal Consistency

- **How to integrate optical flow into colorizing video**

method 2: use output as input

[vid2vid: Video-to-Video Synthesis (2018, Wang)]



Final Generate Image
 $= (1 - mask_t) \odot gen_{flow_{t-1}}(gen_{t-1}) + mask_t \odot gen_{t-1}$

III. Project I: Revisit

G. Temporal Consistency

- **How to integrate optical flow into colorizing video**

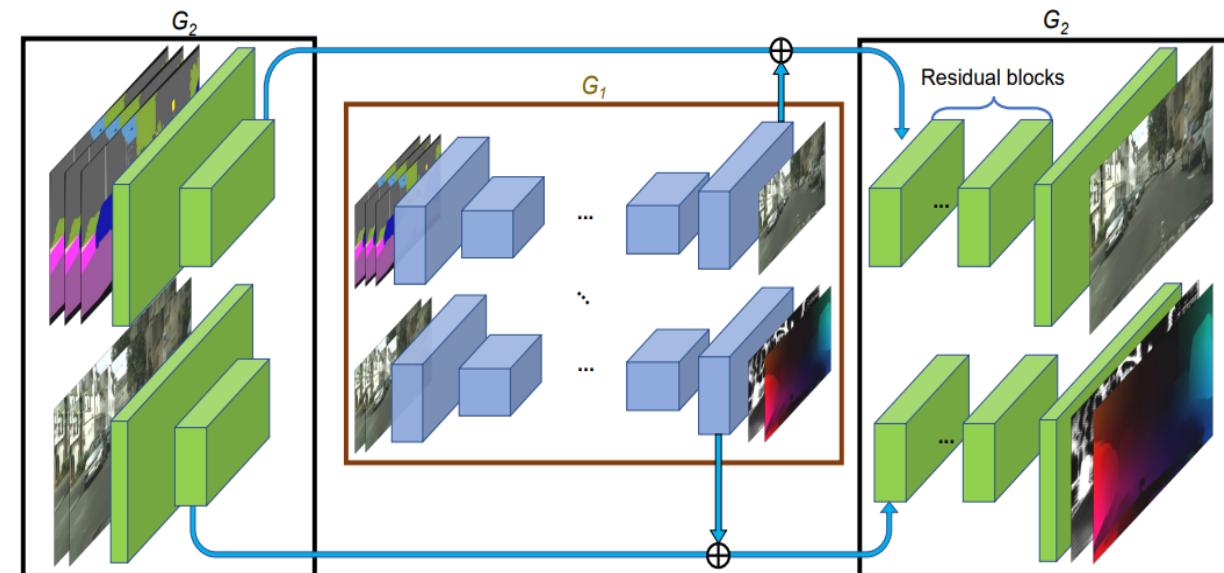
method 2: use output as input

[vid2vid: Video-to-Video Synthesis (2018, Wang)]

Generalized tricks:

1. Multiscale Generator
2. Multiscale Discriminator
3. Inner Loss of Discriminator

[Just like SPADE]



III. Project I: Revisit

G. Temporal Consistency

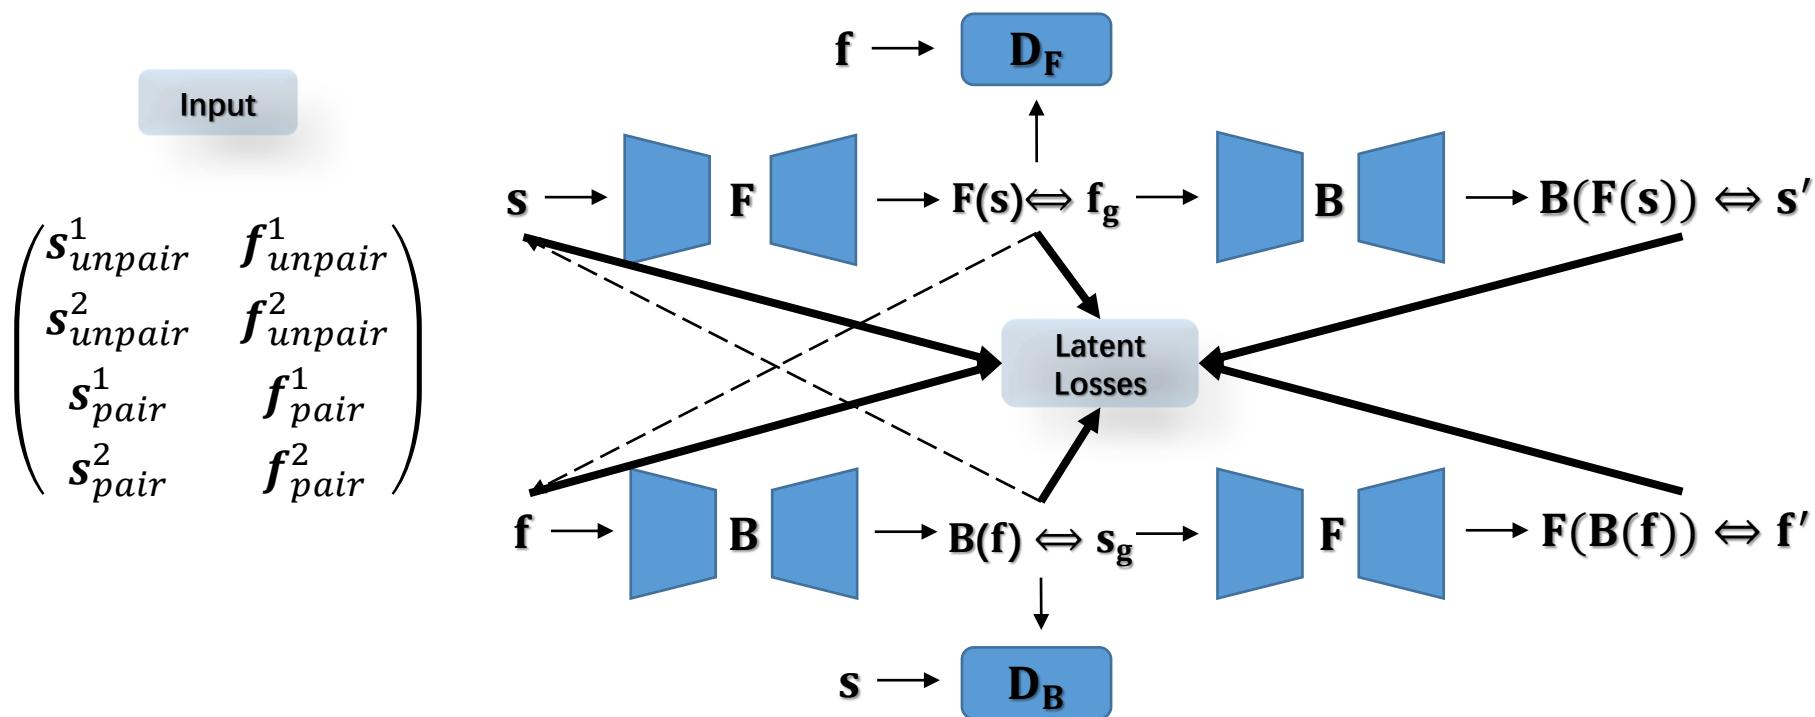
- **How to integrate optical flow into colorizing video**

method 3: other related methods: [link](#)

IV. Project II: Revisit

H. Pair with Unpair

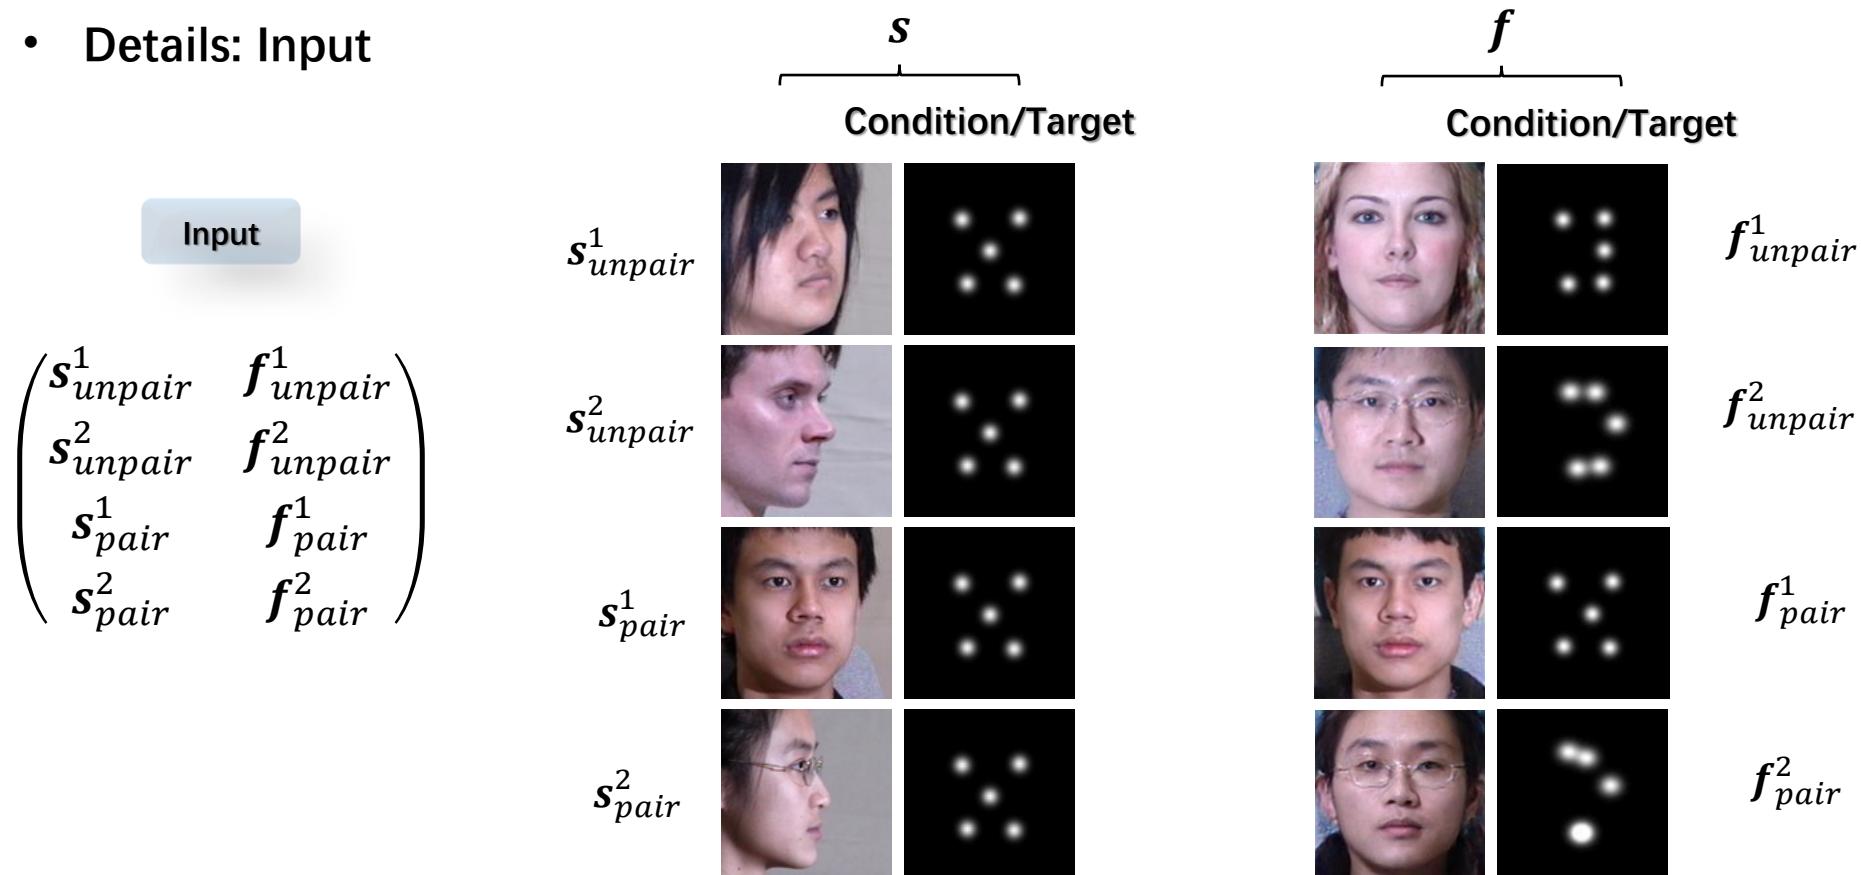
- Structure: Outline



IV. Project II: Revisit

H. Pair with Unpair

- Details: Input



IV. Project II: Revisit

H. Pair with Unpair

- Losses

1. GAN Losses

For $\mathbf{F}: s \rightarrow f$

$$\min_{\mathbf{D}_F} \max_{\mathbf{F}} \mathcal{L}_{GAN}(F, D_F, s, f)$$

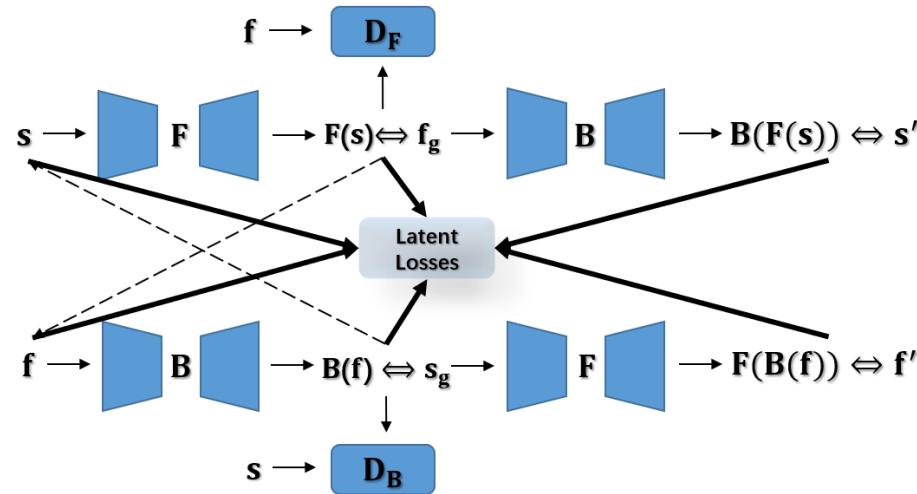
For $\mathbf{B}: f \rightarrow s$

$$\min_{\mathbf{D}_B} \max_{\mathbf{B}} \mathcal{L}_{GAN}(B, D_B, f, s)$$

$$\begin{aligned} \mathcal{L}_{GAN}(F, D_F, s, f) = & \mathbb{E}_{f \sim p_{data}(f)} [\log D_F(f)] \\ & + \mathbb{E}_{s \sim p_{data}(s)} [\log(1 - D_F(F(s)))] \end{aligned}$$

$$\begin{aligned} \mathcal{L}_{GAN}(B, D_B, f, s) = & \mathbb{E}_{s \sim p_{data}(s)} [\log D_B(s)] \\ & + \mathbb{E}_{f \sim p_{data}(f)} [\log(1 - D_B(B(f)))] \end{aligned}$$

Implementation was completed by WGAN-GP



2. Cycle Loss

$$\mathcal{L}_{cyc}(F, B) = \mathbb{E}_{s \sim p_{data}(s)} [\| s' - s \|_1] + \mathbb{E}_{f \sim p_{data}(f)} [\| f' - f \|_1]$$

3. Identity Loss

4. tv (total variation) Loss

$$\mathcal{L}_{tv}(F, B) = f_g^2_{\text{grad}} + f'^2_{\text{grad}} + s_g^2_{\text{grad}} + s'^2_{\text{grad}}$$

IV. Project II: Revisit

H. Pair with Unpair

- Losses

5. Latent Losses

For pair data:

$$\text{inner loss} \sum_l \left(\frac{\mathbf{f} / \mathbf{f}_g}{\mathbf{s}} \rightarrow \mathbf{D}_{fs} + \frac{\mathbf{s} / \mathbf{s}_g}{\mathbf{f}} \rightarrow \mathbf{D}_{sf} \right) + \text{adv loss} + \text{pixel loss} + \text{perceptual loss}$$

+
 $\frac{\mathbf{f}}{\mathbf{f}_g} \rightarrow \text{Light CNN}$

For unpair data:

$$\text{inner loss} \sum_l \left(\frac{\mathbf{s} / \mathbf{s}'}{\mathbf{f}_g} \rightarrow \mathbf{D}_{sf} + \frac{\mathbf{f} / \mathbf{f}'}{\mathbf{s}_g} \rightarrow \mathbf{D}_{fs} \right)$$

